

# Lecture 3: Kernelization

**Making linear models non-linear**

Joaquin Vanschoren

# Feature Maps

- Linear models:  $\hat{y} = \mathbf{w}\mathbf{x} + w_0 = \sum_{i=1}^p w_i x_i + w_0 = w_0 + w_1 x_1 + \dots + w_p x_p$
- When we cannot fit the data well, we can add non-linear transformations of the features
- Feature map (or *basis expansion*)  $\phi : X \rightarrow \mathbb{R}^d$

$$y = \mathbf{w}^T \mathbf{x} \rightarrow y = \mathbf{w}^T \phi(\mathbf{x})$$

- E.g. Polynomial feature map: all polynomials up to degree  $d$  and all products


$$[1, x_1, \dots, x_p] \xrightarrow{\phi} [1, x_1, \dots, x_p, x_1^2, \dots, x_p^2, \dots, x_p^d, x_1 x_2, \dots, x_{p-1} x_p]$$

- Example with  $p = 1, d = 3$ :

$$y = w_0 + w_1 x_1 \xrightarrow{\phi} y = w_0 + w_1 x_1 + w_2 x_1^2 + w_3 x_1^3$$

# Ridge regression example


Weights: [0.418]



- Add all polynomials  $x^d$  up to degree 10 and fit again:
  - e.g. use sklearn `PolynomialFeatures`

	$x_0$	$x_0^2$	$x_0^3$	$x_0^4$	$x_0^5$	$x_0^6$	$x_0^7$	$x_0^8$
0	-0.752759	0.566647	-0.426548	0.321088	-0.241702	0.181944	-0.136960	0.103098
1	2.704286	7.313162	19.776880	53.482337	144.631526	391.124988	1057.713767	2860.360362
2	1.391964	1.937563	2.697017	3.754150	5.225640	7.273901	10.125005	14.093639
3	0.591951	0.350406	0.207423	0.122784	0.072682	0.043024	0.025468	0.015076
4	-2.063888	4.259634	-8.791409	18.144485	-37.448187	77.288869	-159.515582	329.222321

Weights: [ 0.643 0.297 -0.69 -0.264 0.41 0.096 -0.076 -0.014 0.004 0.001]



How expensive is this?

- You may need MANY dimensions to fit the data
  - Memory and computational cost
  - More weights to learn, more likely overfitting
- Ridge has a closed-form solution which we can compute with linear algebra:


$$w^* = (X^T X + \alpha I)^{-1} X^T Y$$

- Since  $X$  has  $n$  rows (examples), and  $d$  columns (features),  $X^T X$  has dimensionality  $d \times d$
- Hence Ridge is quadratic in the number of features,  $\mathcal{O}(d^2 n)$
- After the feature map  $\Phi$ , we get


$$w^* = (\Phi(X)^T \Phi(X) + \alpha I)^{-1} \Phi(X)^T Y$$

- Since  $\Phi$  increases  $d$  a lot,  $\Phi(X)^T \Phi(X)$  becomes *huge*


# Linear SVM example (classification)



We can add a new feature by taking the squares of feature1 values




Now we can fit a linear model





As a function of the original features, the decision boundary is now a polynomial as well

$$y = w_0 + w_1x_1 + w_2x_2 + w_3x_2^2 > 0$$



# The kernel trick

- Computations in explicit, high-dimensional feature maps are *expensive*
- For *some* feature maps, we can, however, compute *distances* between *points* cheaply
  - Without explicitly constructing the high-dimensional space at all
- Example: *quadratic* feature map for  $\mathbf{x} = (x_1, \dots, x_p)$ :

$$\Phi(\mathbf{x}) = (x_1, \dots, x_p, x_1^2, \dots, x_p^2, \sqrt{2}x_1x_2, \dots, \sqrt{2}x_{p-1}x_p)$$

- A *kernel function* exists for this feature map to compute dot products


$$k_{quad}(\mathbf{x}_i, \mathbf{x}_j) = \Phi(\mathbf{x}_i) \cdot \Phi(\mathbf{x}_j) = \mathbf{x}_i \cdot \mathbf{x}_j + (\mathbf{x}_i \cdot \mathbf{x}_j)^2$$

- Skip computation of  $\Phi(x_i)$  and  $\Phi(x_j)$  and compute  $k(x_i, x_j)$  directly


# Kernelization

- Kernel  $k$  corresponding to a feature map  $\Phi$ :  $k(\mathbf{x}_i, \mathbf{x}_j) = \Phi(\mathbf{x}_i) \cdot \Phi(\mathbf{x}_j)$
- Computes dot product between  $x_i, x_j$  in a high-dimensional space  $\mathcal{H}$ 
  - Kernels are sometimes called *generalized dot products*
  - $\mathcal{H}$  is called the *reproducing kernel Hilbert space* (RKHS)
- The dot product is a measure of the *similarity* between  $x_i, x_j$ 
  - Hence, a kernel can be seen as a similarity measure for high-dimensional spaces
- If we have a loss function based on dot products  $\mathbf{x}_i \cdot \mathbf{x}_j$  it can be *kernelized*
  - Simply replace the dot products with  $k(\mathbf{x}_i, \mathbf{x}_j)$

Low-dimensional space



High-dimensional space (RKHS)



$$k(x_1, x_2) = \Phi(x_1) \cdot \Phi(x_2)$$

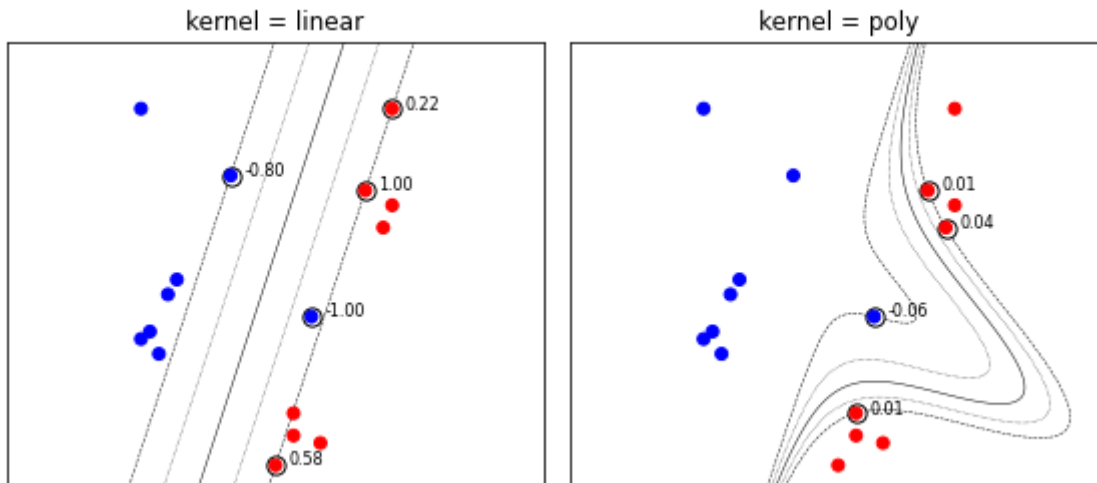
# Example: SVMs

- Linear SVMs (dual form, for  $l$  support vectors with dual coefficients  $a_i$  and classes  $y_i$ ):

$$\mathcal{L}_{Dual}(a_i) = \sum_{i=1}^l a_i - \frac{1}{2} \sum_{i,j=1}^l a_i a_j y_i y_j (\mathbf{x}_i \cdot \mathbf{x}_j)$$

- Kernelized SVM, using any existing kernel  $k$  we want:

$$\mathcal{L}_{Dual}(a_i, k) = \sum_{i=1}^l a_i - \frac{1}{2} \sum_{i,j=1}^l a_i a_j y_i y_j k(\mathbf{x}_i, \mathbf{x}_j)$$



## Which kernels exist?


- A (Mercer) kernel is any function  $k : X \times X \rightarrow \mathbb{R}$  with these properties:
  - Symmetry:  $k(\mathbf{x}_1, \mathbf{x}_2) = k(\mathbf{x}_2, \mathbf{x}_1) \quad \forall \mathbf{x}_1, \mathbf{x}_2 \in X$
  - Positive definite: the kernel matrix  $K$  is positive semi-definite
    - Intuitively,  $k(\mathbf{x}_1, \mathbf{x}_2) \geq 0$
- The kernel matrix (or Gram matrix) for  $n$  points of  $x_1, \dots, x_n \in X$  is defined as:

$$K = XX^T = \begin{bmatrix} k(\mathbf{x}_1, \mathbf{x}_1) & \dots & k(\mathbf{x}_1, \mathbf{x}_n) \\ \vdots & \ddots & \vdots \\ k(\mathbf{x}_n, \mathbf{x}_1) & \dots & k(\mathbf{x}_n, \mathbf{x}_n) \end{bmatrix}$$


- Once computed ( $\mathcal{O}(n^2)$ ), simply lookup  $k(\mathbf{x}_1, \mathbf{x}_2)$  for any two points
- In practice, you can either supply a kernel function or precompute the kernel matrix

# Linear kernel

- Input space is same as output space:  $X = \mathcal{H} = \mathbb{R}^d$
- Feature map  $\Phi(\mathbf{x}) = \mathbf{x}$
- Kernel:  $k_{linear}(\mathbf{x}_i, \mathbf{x}_j) = \mathbf{x}_i \cdot \mathbf{x}_j$
- Geometrically, the dot product is the *projection* of  $\mathbf{x}_j$  on hyperplane defined by  $\mathbf{x}_i$ 
  - Becomes larger if  $\mathbf{x}_i$  and  $\mathbf{x}_j$  are in the same 'direction'




- Linear kernel between point  $(0,1)$  and another unit vector an angle  $\alpha$  (in radians)
  - Points with similar angles are deemed similar



# Polynomial kernel

- If  $k_1, k_2$  are kernels, then  $\lambda \cdot k_1$  ( $\lambda \geq 0$ ),  $k_1 + k_2$ , and  $k_1 \cdot k_2$  are also kernels
- The **polynomial kernel** (for degree  $d \in \mathbb{N}$ ) reproduces the polynomial feature map
  - $\gamma$  is a scaling hyperparameter (default  $\frac{1}{p}$ )
  - $c_0$  is a hyperparameter (default 1) to trade off influence of higher-order terms

$$k_{poly}(\mathbf{x}_1, \mathbf{x}_2) = (\gamma(\mathbf{x}_1 \cdot \mathbf{x}_2) + c_0)^d$$






# RBF (Gaussian) kernel

- The *Radial Basis Function* (RBF) feature map builds the Taylor series expansion of  $e^x$


$$\Phi(x) = e^{-x^2/2\gamma^2} \left[ 1, \sqrt{\frac{1}{1!\gamma^2}} x, \sqrt{\frac{1}{2!\gamma^4}} x^2, \sqrt{\frac{1}{3!\gamma^6}} x^3, \dots \right]^T$$

- RBF (or *Gaussian*) kernel with *kernel width*  $\gamma \geq 0$ :

$$k_{RBF}(\mathbf{x}_1, \mathbf{x}_2) = \exp(-\gamma \|\mathbf{x}_1 - \mathbf{x}_2\|^2)$$




- The RBF kernel  $k_{RBF}(\mathbf{x}_1, \mathbf{x}_2) = \exp(-\gamma \|\mathbf{x}_1 - \mathbf{x}_2\|^2)$  does not use a dot product
  - It only considers the distance between  $\mathbf{x}_1$  and  $\mathbf{x}_2$
  - It's a *local kernel* : every data point only influences data points nearby
    - linear and polynomial kernels are *global* : every point affects the whole space
  - Similarity depends on closeness of points and kernel width
    - value goes up for closer points and wider kernels (larger overlap)



# Kernelized SVMs in practice

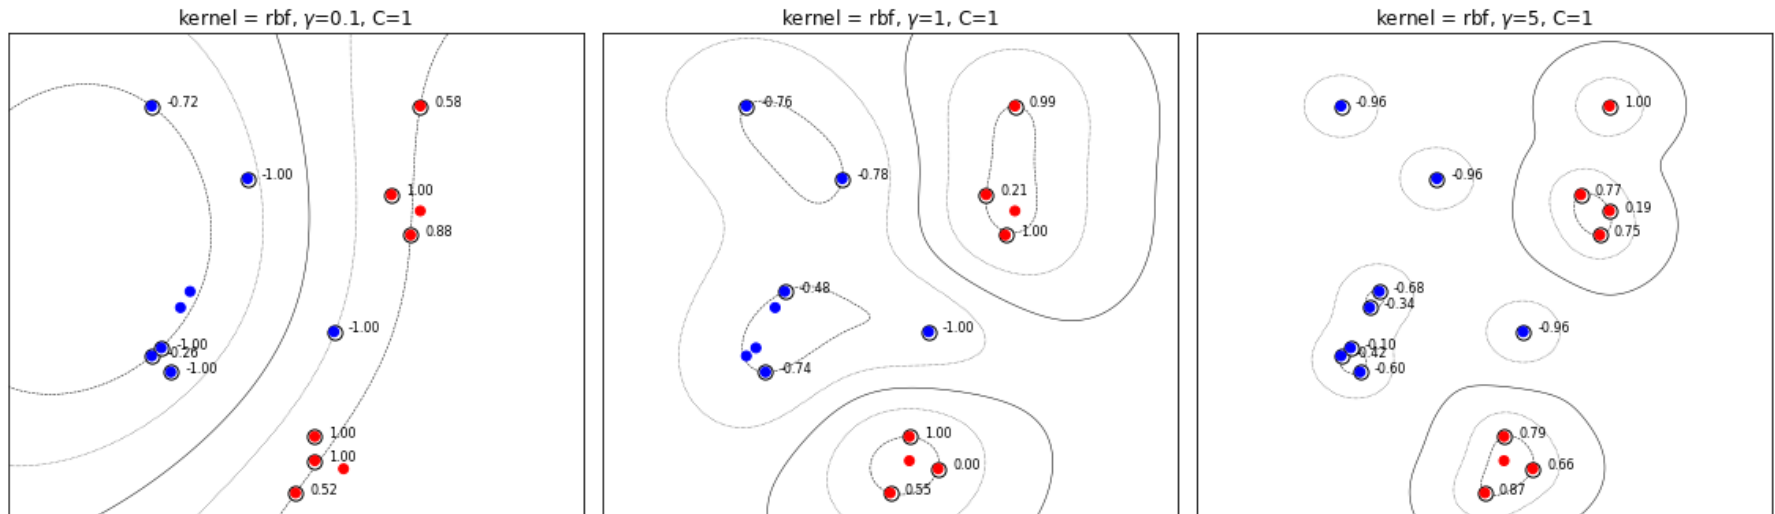
- You can use SVMs with any kernel to learn non-linear decision boundaries



# SVM with RBF kernel

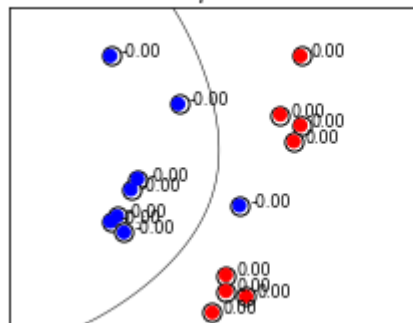

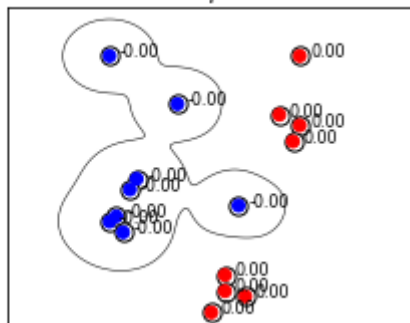






- Every support vector *locally* influences predictions, according to kernel width ( $\gamma$ )
- The prediction for test point  $\mathbf{u}$ : sum of the remaining influence of each support vector

$$\blacksquare f(x) = \sum_{i=1}^l a_i y_i k(\mathbf{x}_i, \mathbf{u})$$

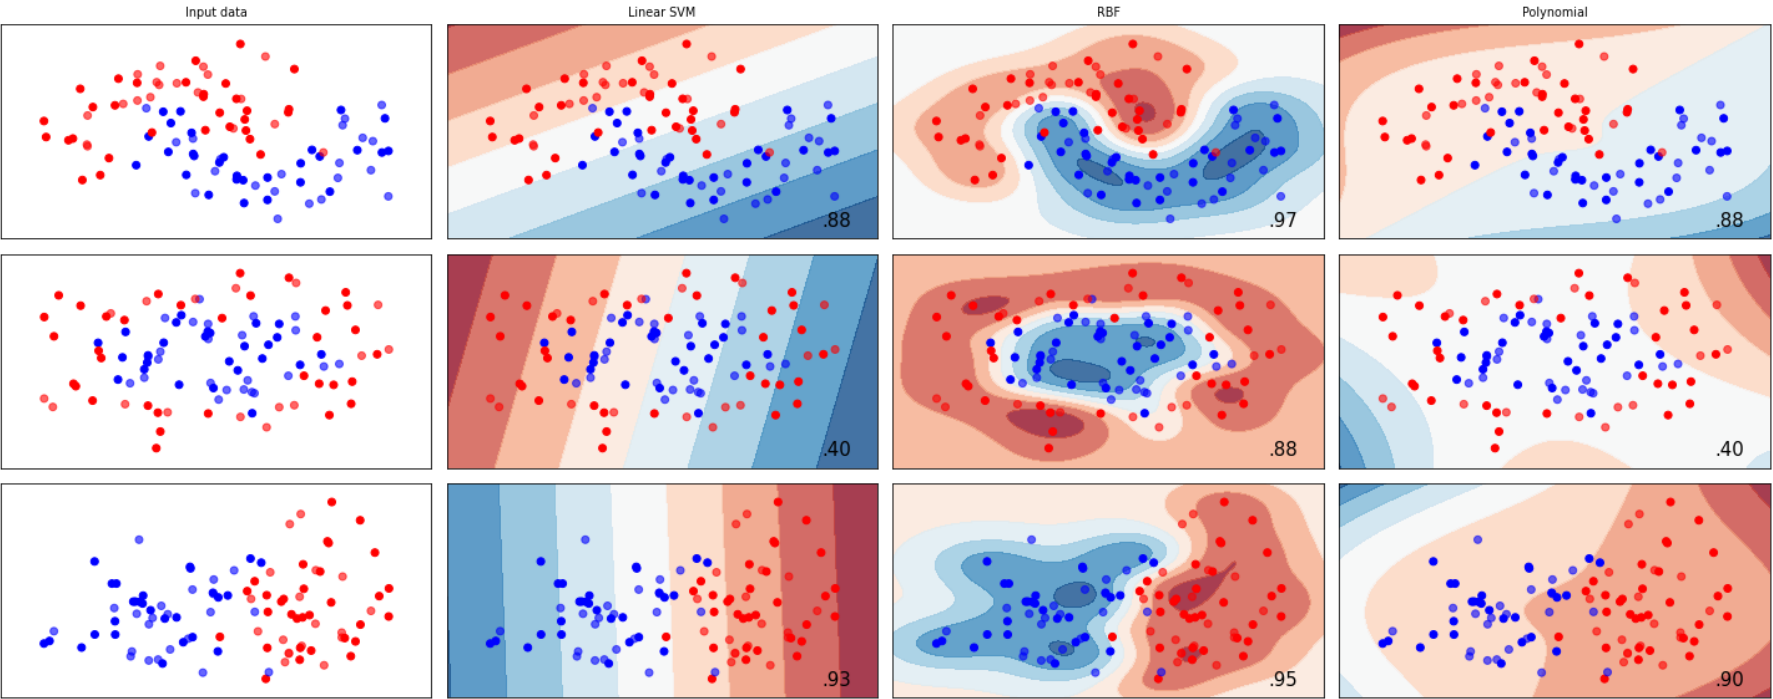


## Tuning RBF SVMs

- gamma (kernel width)
  - high values cause narrow Gaussians, more support vectors, overfitting
  - low values cause wide Gaussians, underfitting
- C (cost of margin violations)
  - high values punish margin violations, cause narrow margins, overfitting
  - low values cause wider margins, more support vectors, underfitting

kernel = rbf,  $\gamma=0.1$ ,  $C=0.001$ kernel = rbf,  $\gamma=1$ ,  $C=0.001$ kernel = rbf,  $\gamma=5$ ,  $C=0.001$ kernel = rbf,  $\gamma=0.1$ ,  $C=1$ kernel = rbf,  $\gamma=1$ ,  $C=1$ kernel = rbf,  $\gamma=5$ ,  $C=1$ kernel = rbf,  $\gamma=0.1$ ,  $C=100$ kernel = rbf,  $\gamma=1$ ,  $C=100$ kernel = rbf,  $\gamma=5$ ,  $C=100$ 

Kernel overview



## SVMs in practice

- C and gamma *always* need to be tuned
  - Interacting regularizers. Find a good C, then finetune gamma
- SVMs expect all features to be approximately on the same scale
  - Data needs to be scaled beforehand
- Allow to learn complex decision boundaries, even with few features
  - Work well on both low- and high dimensional data
  - Especially good at small, high-dimensional data
- Hard to inspect, although support vectors can be inspected
- In sklearn, you can use `SVC` for classification with a range of kernels
  - `SVR` for regression



## Other kernels

- There are many more possible kernels
- If no kernel function exists, we can still *precompute* the kernel matrix
  - All you need is some similarity measure, and you can use SVMs
- Text kernels:
  - Word kernels: build a bag-of-words representation of the text (e.g. TFIDF)
    - Kernel is the inner product between these vectors
  - Subsequence kernels: sequences are similar if they share many sub-sequences
    - Build a kernel matrix based on pairwise similarities
- Graph kernels: Same idea (e.g. find common subgraphs to measure similarity)
- These days, deep learning embeddings are more frequently used

# The Representer Theorem

- We can kernelize many other loss functions as well
- The *Representer Theorem* states that if we have a loss function  $\mathcal{L}'$  with
  - $\mathcal{L}$  an *arbitrary* loss function using some function  $f$  of the inputs  $\mathbf{x}$
  - $\mathcal{R}$  a (non-decreasing) regularization score (e.g. L1 or L2) and constant  $\lambda$

$$\mathcal{L}'(\mathbf{w}) = \mathcal{L}(y, f(\mathbf{x})) + \lambda \mathcal{R}(\|\mathbf{w}\|)$$

- Then the weights  $\mathbf{w}$  can be described as a linear combination of the training samples:

$$\mathbf{w} = \sum_{i=1}^n a_i y_i f(\mathbf{x}_i)$$

- Note that this is exactly what we found for SVMs:  $\mathbf{w} = \sum_{i=1}^l a_i y_i \mathbf{x}_i$
- Hence, we can also kernelize Ridge regression, Logistic regression, Perceptrons, Support Vector Regression, ...

# Kernelized Ridge regression

- The linear Ridge regression loss (with  $\mathbf{x}_0 = 1$ ):

$$\mathcal{L}_{Ridge}(\mathbf{w}) = \sum_{i=0}^n (y_i - \mathbf{w}\mathbf{x}_i)^2 + \lambda \|\mathbf{w}\|^2$$

- Filling in  $\mathbf{w} = \sum_{i=1}^n \alpha_i y_i \mathbf{x}_i$  yields the dual formulation:


$$\mathcal{L}_{Ridge}(\mathbf{w}) = \sum_{i=1}^n \left( y_i - \sum_{j=1}^n \alpha_j y_j \mathbf{x}_i \mathbf{x}_j \right)^2 + \lambda \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j \mathbf{x}_i \mathbf{x}_j$$

- Generalize  $\mathbf{x}_i \cdot \mathbf{x}_j$  to  $k(\mathbf{x}_i, \mathbf{x}_j)$


$$\mathcal{L}_{KernelRidge}(\alpha, k) = \sum_{i=1}^n \left( y_i - \sum_{j=1}^n \alpha_j y_j k(\mathbf{x}_i, \mathbf{x}_j) \right)^2 + \lambda \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j k(\mathbf{x}_i, \mathbf{x}_j)$$

## Example of kernelized Ridge

- Prediction (red) is now a linear combination of kernels (blue):  $y = \sum_{j=1}^n \alpha_j y_j k(\mathbf{x}, \mathbf{x}_j)$
- We learn a dual coefficient for each point



- Fitting our regression data with `KernelRidge`

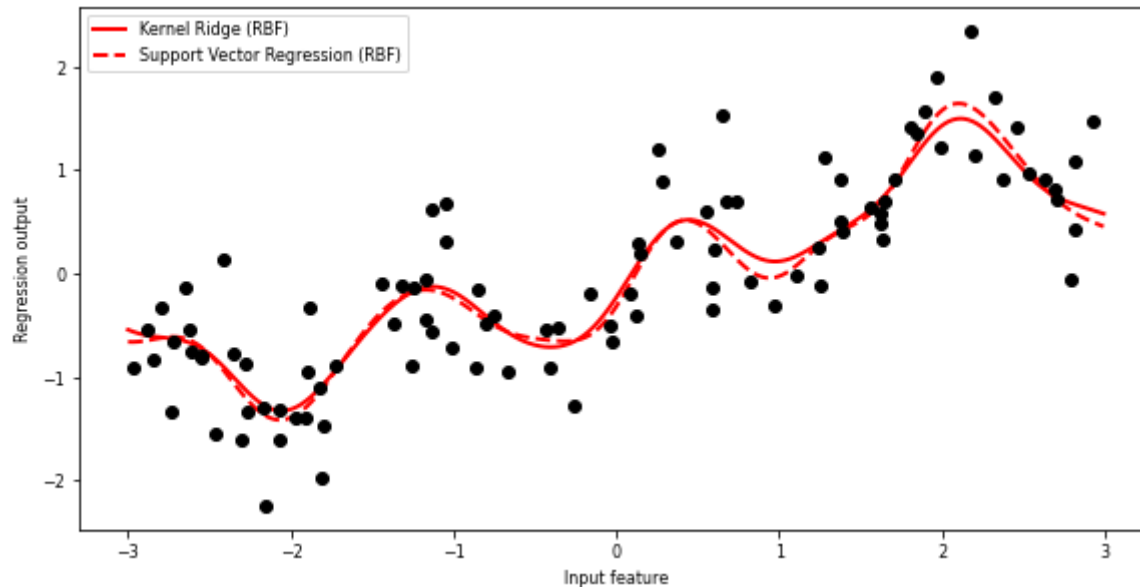


## Other kernelized methods

- Same procedure can be done for logistic regression
- For perceptrons,  $\alpha \rightarrow \alpha + 1$  after every misclassification

$$\mathcal{L}_{DualPerceptron}(x_i, k) = \max(0, y_i \sum_{j=1}^n \alpha_j y_j k(\mathbf{x}_j, \mathbf{x}_i))$$

- Support Vector Regression behaves similarly to Kernel Ridge



# Summary

- Feature maps  $\Phi(x)$  transform features to create a higher-dimensional space
  - Allows learning non-linear functions or boundaries, but very expensive/slow
- For some  $\Phi(x)$ , we can compute dot products without constructing this space
  - Kernel trick:  $k(\mathbf{x}_i, \mathbf{x}_j) = \Phi(\mathbf{x}_i) \cdot \Phi(\mathbf{x}_j)$
  - Kernel  $k$  (generalized dot product) is a measure of similarity between  $\mathbf{x}_i$  and  $\mathbf{x}_j$
- There are many such kernels
  - Polynomial kernel:  $k_{poly}(\mathbf{x}_1, \mathbf{x}_2) = (\gamma(\mathbf{x}_1 \cdot \mathbf{x}_2) + c_0)^d$
  - RBF (Gaussian) kernel:  $k_{RBF}(\mathbf{x}_1, \mathbf{x}_2) = \exp(-\gamma \|\mathbf{x}_1 - \mathbf{x}_2\|^2)$
  - A kernel matrix can be precomputed using any similarity measure (e.g. for text, graphs,...)
- Any loss function where inputs appear only as dot products can be kernelized
  - E.g. Linear SVMs: simply replace the dot product with a kernel of choice
- The Representer theorem states which *other* loss functions can also be kernelized and how
  - Ridge regression, Logistic regression, Perceptrons,...