


# Neural Networks

# Mathematical Foundations

- A first example
- Tensors and tensor operations
- Backpropagation and gradient descent

## A first example: classifying digits

- This example is meant to introduce the main concepts. We'll cover them in more detail later.
- MNIST dataset contains 28x28 pixel images of handwritten digits (0-9)
- The goal is to classify each image as one of the possible digits
- We **reshape** the data to a 70000x28x28 **tensor** (n-dimensional matrix)  
`X = X.reshape(70000, 28, 28)`
- Traditional holdout uses the last 10,000 images for testing



Training set: (60000, 28, 28)


Test set: (10000, 28, 28)

Note: this is also one of the datasets that comes included with Keras:


```
from keras.datasets import mnist  
(train_images, train_labels), (test_images, test_labels) = mnist.load_data()  
ta()
```

## Neural networks

- The core building block of a neural network is the *layer*
- You can think of it as a *filter* for the data
  - Data goes in, and comes out in a more useful form
- Layers extract new *representations* of the data
- *Deep learning* models contain many such layers
  - They progressively *distill* (refine) the data




- The pixel values are fed to individual *nodes* of the *input layer* (yellow)
- The data then passes through one or more *hidden layers* (blue)
  - One type of layer is the *dense* or *fully connected* layer
  - Every node is connected to all nodes in the previous and subsequent layers
- The *output layer* has a node for every possible outcome (digits 0-9) (red)
  - I.e. The first node returns the probability that the input image represents a '0'




## The perceptron

- In its simplest form, each node outputs a weighted sum of the inputs:  
$$y = \sum_i x_i w_i + b$$
- It needs to learn the optimal set of weights to produce the right output
  - The *bias*  $b$  is modelled as the weight of an extra input that is always '1' (see later)
- This is exactly the same as a linear model. It can only learn linear decision boundaries.
  - Even a deep neural net of perceptrons can only learn a linear model



## Activation functions


- To learn a non-linear model, each hidden node has to output a non-linear *activation function*  $f$  on the weighted sum of the inputs:  
$$h(x) = f(W_1x + b_1)$$
- Likewise, the output nodes use an activation function  $g$  on the weighted outputs of the previous layer:  $o(x) = g(W_2h(x) + b_2)$





## Activation functions

- For hidden nodes, popular choices are the *rectified linear unit* (ReLU) and *tanh*
  - There are many others. We'll come back to this soon!
  - ReLU is very cheap to compute, speeds up training
- For classification, we use *softmax* (or sigmoid)
  - Transforms the input into a probability for (each specific outcome)
  - This is exactly what we used for logistic regression!



We can now build a simple neural network for MNIST:

- One dense hidden ReLU layer with 512 nodes
  - Input from a 28x28 matrix
- Output softmax layer with 10 nodes

```
from keras import models  
from keras import layers
```

```
network = models.Sequential()  
network.add(layers.Dense(512, activation='relu', input_shape=(28 * 28  
,)))  
network.add(layers.Dense(10, activation='softmax'))
```

'Visualize' the model using `summary( )`

- Also shows the number of model parameters (weights) that need to be learned

Layer (type)	Output Shape	Param #
dense_3 (Dense)	(None, 512)	401920
dense_4 (Dense)	(None, 10)	5130
Total params: 407,050		
Trainable params: 407,050		
Non-trainable params: 0		

## Compilation

We still need to specify how we want the network to be trained:

- **Loss function:** The objective function used to measure how well the model is doing, and steer itself in the right direction
  - e.g. Cross Entropy (*negative log likelihood* or *log loss*) for classification
- **Optimizer:** How to optimize the model weights in every iteration.
  - usually a variant of stochastic gradient descent (<http://runder.io/optimizing-gradient-descent/index.html#momentum>).
- **Metrics** to monitor performance during training and testing.
  - e.g. accuracy

```
network.compile(optimizer='rmsprop',  
                loss='categorical_crossentropy',  
                metrics=['accuracy'])
```

## Preprocessing

- Neural networks are sensitive to scaling, so always scale the inputs
- The network expects the data in shape  $(n, 28 * 28)$ , so we also need to reshape
- We also need to categorically encode the labels
  - e.g. '4' becomes  $[0,0,0,0,1,0,0,0,0,0]$

```
from keras.utils import to_categorical  
y_train = to_categorical(y_train)  
y_test = to_categorical(y_test)
```

## Training

Training (fitting) is done by **stochastic gradient descent** (SGD).

- Optimizes the model parameters (weights)
- We'll come back to this soon

```
Epoch 1/5
60000/60000 [=====] - 5s 79us/step - loss: 0.2536
- acc: 0.9265
Epoch 2/5
60000/60000 [=====] - 4s 75us/step - loss: 0.1032
- acc: 0.9699
Epoch 3/5
60000/60000 [=====] - 5s 79us/step - loss: 0.0671
- acc: 0.9795
Epoch 4/5
60000/60000 [=====] - 4s 74us/step - loss: 0.0491
- acc: 0.9854
Epoch 5/5
60000/60000 [=====] - 4s 72us/step - loss: 0.0362
- acc: 0.9889
```

## Prediction

We can now call `predict` or `predict_proba` to generate predictions

```
Prediction: [0.          0.0000002 0.          0.000001  0.9863701 0.000019
0.
0.0000008 0.0005651 0.0130437]
Label: [0. 0. 0. 0. 1. 0. 0. 0. 0. 0.]
```

## Evaluation

Evaluate the trained model on the entire test set

```
10000/10000 [=====] - 1s 57us/step  
Test accuracy: 0.9793
```



## Overfitting


- Our test set accuracy is quite a bit lower than our training set accuracy
- We've already seen many choices (moving pieces) that can still be optimized:
  - Number of layers
  - Number of nodes per layer
  - Activation functions
  - Loss function (and hyperparameters)
  - SGD optimizer (and hyperparameters)
  - Batch size
  - Number of epochs

# **Tensors and tensor operations**

Representing data and learning better representations

## Tensors

- A *tensor* is simply an n-dimensional array (with n axes)
  - 2D tensor: matrix (samples, features)
  - 3D tensor: grayscale images (samples, height, width)
    - or time series (samples, timesteps, features)
  - 4D tensor: color images (samples, height, width, channels)
  - 5D tensor: video (amples, frames, height, width, channels)



## Tensor operations

The operations that neural network layers perform on the data can be reduced to a handful of tensor operations.

```
keras.layers.Dense(512, activation='relu')
```

can be interpreted as a function

```
y = relu(dot(W, x) + b)
```

- takes a 2D tensor  $x$  and returns a new 2D tensor  $y$
- uses a 2D weight tensor  $W$  and a bias vector  $b$
- performs a dot product, addition, and  $relu(x) = \max(x, 0)$

## Element-wise operations

ReLU and addition are element-wise operations. Since numpy arrays support element-wise operations natively, these are simply:

```
def relu(x):  
    return np.maximum(x, 0.)
```

```
def add(x, y):  
    return x + y
```

Note: if  $y$  has a lower dimension than  $x$ , it will be *broadcasted*: axes are added to match the dimensionality, and  $y$  is repeated along the new axes

```
>>> np.array([[1,2],[3,4]]) + np.array([10,20])  
array([[11, 22],  
       [13, 24]])
```


## Tensor dot

The dot product  $x \cdot y$  of two tensors can also be done easily with numpy:

```
z = np.dot(x, y)
```


where

```
z[i,j] = x[i,:] * y[:,j]
```



## Geometric interpretation

- Dot products and additions change how data points relate to each other
- We aim to find a transformation of the data points so that it becomes easy to:
  - separate the classes (classification)
  - learn a simple function (regression)



# Gradient-based optimization

- We saw that a layer performs an operation like:

$$y = \text{relu}(\text{dot}(W, x) + b)$$

- How to find good values for  $W$  and  $b$  so that the data is transformed to a useful representation?
- Start with a random initialization, then loop:
  1. Draw a batch of training data  $x$
  2. *Forward pass*: run the network on  $x$  to yield  $y_{pred}$  (tensor operations)
  3. Compute the loss (mismatch between  $y_{pred}$  and  $y$ )
  4. Update  $W, b$  in a way that slightly reduces the loss (OK, but how?)




## Update rule

Naive approach (expensive):

- Choose one weight  $w_{i,j}$  to optimize, freeze the others
- Run the network (twice) with  $w_{i,j} - \epsilon$  and  $w_{i,j} + \epsilon$
- Compute the losses given current batch  $x$
- Keep the one that reduces the loss most, then repeat


Better:



- Choose a loss function  $f$  that is *differentiable*
  - Also all underlying tensor operations need to be differentiable
- Then we can compute the derivative  $\frac{\partial f(x, w_{i,j})}{\partial w_{i,j}} = a$
- So that  $f(x, w_{i,j} + \epsilon) = y + a * \epsilon$
- We can now estimate better weights without recomputing  $f$

## Gradients

- A *gradient* is the generalization of a derivate to n-dimensional inputs
  - Approximates the *curvature* of the loss function  $f(x, W)$  around a given point  $W$
- Update: if  $f$  is differentiable, then  $W_1 = W_0 - \frac{\partial f(W_0)}{\partial W} * step$ 
  - step is a small scaling factor
  - Go against the curvature to a lower place on the curve
- Now repeat with a new batch of data  $x$



## Stochastic gradient descent (SGD)



Mini-batch SGD:

1. Draw a batch of *batch\_size* training data  $x$  and  $y$
2. *Forward pass*: run the network on  $x$  to yield  $y_{pred}$  (tensor operations)
3. Compute the loss  $L$  (mismatch between  $y_{pred}$  and  $y$ )
4. *Backward pass*: Compute the gradient of the loss with regard to  $W$
5. Update  $W$ :  $W_{i+1} = W_i - \frac{\partial L(x, W_i)}{\partial W} * step$

Repeat until  $n$  passes (epochs) are made through through the entire training set.


## SGD Variants:

- Batch Gradient Descent: compute gradient on entire training set
  - More accurate gradients, but more expensive
- True Stochastic Gradient Descent: repeat for each individual data points (noisy)
- Minibatch SGD strikes a balance between the two (given the right batch size)



## SGD: many more variants

- With SGD, it is quite easy to get stuck in a local minimum
- Learning rate decay: start with a big step size and then decrease
- Momentum: do a larger update if previous update has large loss improvement
  - Like a ball that gains speed if it goes down steeply
- Adaptive step size for each  $W_i$ : adam, Adagrad, ...
  - See <http://ruder.io/optimizing-gradient-descent/index.html> (<http://ruder.io/optimizing-gradient-descent/index.html>)
- Some intuitions say that in high-dimensional spaces, most local minima are near the global minimum



## Backpropagation

- In practice, a neural network function consist of many tensor operations chained together
  - e.g.  $f(W1, W2, W3) = a(W1, b(W2, c(W3)))$
- As long as each tensor operation is differentiable, we can still compute the gradient thanks to the chain rule:

$$f(g(x)) = f'(g(x)) * g'(x)$$

- We can let the gradient *backpropagate* through the layers
- So, if we have a hidden node  $h(x) = f(W_1x + b_1)$ ,  $net(x) = W_1x + b_1$ , and output node  $o(x) = g(W_2h(x) + b_2)$

$$\frac{\partial o(\mathbf{x})}{\partial W_1} = \underbrace{\frac{\partial o(\mathbf{x})}{\partial h(\mathbf{x})}}_{\text{backpropagation of gradient of layer above.}} \underbrace{\frac{\partial h(\mathbf{x})}{\partial net(\mathbf{x})}}_{\text{Gradient of Non-linearity f}} \underbrace{\frac{\partial net(\mathbf{x})}{\partial W_1}}_{\text{Input to 1st layer x}}$$

## Symbolic and automatic differentiation

Symbolic differentiation: given a chain of operations with a known derivative, we can compute a *gradient function* for the chain

- Decomposes functions into simpler functions via the chain rule
- We can call the gradient function to get the gradient value for every model parameter

Automatic differentiation: evaluate the derivative of a function numerically for faster calculation

Modern tools such as TensorFlow do this for you so you don't have to implement backpropagation