

Linear models

Linear models make a prediction using a linear function of the input features.
Can be very powerful for or datasets with many features.

If you have more features than training data points, any target y can be perfectly modeled (on the training set) as a linear function.

Linear models for regression


Prediction formula for input features x . w_i and b are the *model parameters* that need to be learned.

$$\hat{y} = w_0 * x_0 + w_1 * x_1 + \dots + w_p * x_p + b$$

There are many different algorithms, differing in how w and b are learned from the training data.



```
In [2]: mglearn.plots.plot_linear_regression_  
wave()
```

```
w[0]: 0.393906  b: -  
0.031804
```



Linear Regression aka Ordinary Least Squares

- Finds the parameters w and b that minimize the *mean squared error* between predictions and the true regression targets, y , on the training set.
 - MSE: Sum of the squared differences between the predictions and the true values.
- Convex optimization problem with unique closed-form solution (if you have more data points than model parameters w)
- It has no hyperparameters, thus model complexity cannot be controlled.



Linear regression can be found in `sklearn.linear_model`. We'll evaluate it on the Boston Housing dataset.

```
In [3]: from sklearn.model_selection import train_test_split  
from sklearn.linear_model import LinearRegression  
X, y = mglearn.datasets.load_extended_boston()  
  
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)  
lr = LinearRegression().fit(X_train, y_train)
```

```
In [4]: print("Weights (coefficients): {}".format(lr.coef_))
print("Bias (intercept): {}".format(lr.intercept_))
```

```
Weights (coefficients): [ -402.752  -
50.071 -133.317  -12.002  -12.711
28.305  54.492
-51.734  25.26  36.499 -10.1
04 -19.629 -21.368  14.647
2895.054 1510.269 117.995 -26.5
66 31.249 -31.446 45.254
1283.496 -2246.003 222.199 -0.4
66 40.766 -13.436 -19.096
-2.776 -80.971 9.731 5.1
33 -0.788 -7.603 33.672
-11.505 66.267 -17.563 42.9
83 1.277 0.61 57.187
14.082 55.34 -30.348 18.8
12 -13.777 60.979 -12.579
-12.002 -17.698 -34.028 7.1
5 -8.41 16.986 -12.941
-11.806 57.133 -17.581 1.6
96 27.218 -16.745 75.03
-30.272 47.78 -40.541 5.5
04 21.531 25.366 -49.485
28.109 10.469 -71.559 -23.7
4 9.574 -3.788 1.214
-4.72 41.238 -37.702 -2.1
56 -26.296 -33.202 45.932
-23.014 -17.515 -14.085 -20.4
9 36.525 -94.897 143.234
-15.674 -14.973 -28.613 -31.2
52 24.565 -17.805 4.035
1.711 34.474 11.219 1.1
43 3.737 31.385]
Bias (intercept): 31.645174100825688
```

Ridge regression

- Same formula as linear regression
- Adds a penalty term to the least squares sum : $\lambda \sum_i w_i^2$
- Requires that the coefficients (w) are close to zero.
 - Each feature should have as little effect on the outcome as possible
- Regularization: explicitly restrict a model to avoid overfitting.
- Type of L2 regularization: prefers many small weights
 - L1 regularization prefers sparsity: many weights to be 0, others large

Ridge can also be found in `sklearn.linear_model`.

```
In [6]: from sklearn.linear_model import Ridge  
  
ridge = Ridge().fit(X_train, y_train)  
print("Training set score: {:.2f}".format(ridge.  
score(X_train, y_train)))  
print("Test set score: {:.2f}".format(ridge.scor  
e(X_test, y_test)))
```

Training s
et score:
0.89
Test set s
core: 0.75

Test set score is higher and training set score lower: less overfitting!

The strength of the regularization can be controlled with the `alpha` parameter. Default is 1.0.

- Increasing `alpha` forces coefficients to move more toward zero (more regularization)
- Decreasing `alpha` allows the coefficients to be less restricted (less regularization)

```
In [7]: ridge10 = Ridge(alpha=10).fit(X_train, y_train)
print("Training set score: {:.2f}".format(ridge10.score(X_train, y_train)))
print("Test set score: {:.2f}".format(ridge10.score(X_test, y_test)))
```

```
Training
set score: 0.79
Test set
score: 0.64
```

```
In [8]: ridge01 = Ridge(alpha=0.1).fit(X_train, y_train)
print("Training set score: {:.2f}".format(ridge01.score(X_train, y_train)))
print("Test set score: {:.2f}".format(ridge01.score(X_test, y_test)))
```


```
Training
set score: 0.93
Test set
score: 0.77
```


We can plot the weight values for different levels of regularization.

```
In [9]: plt.plot(ridge.coef_, 's', label="Ridge alpha=1")
plt.plot(ridge10.coef_, '^', label="Ridge alpha=10")
plt.plot(ridge01.coef_, 'v', label="Ridge alpha=0.1")

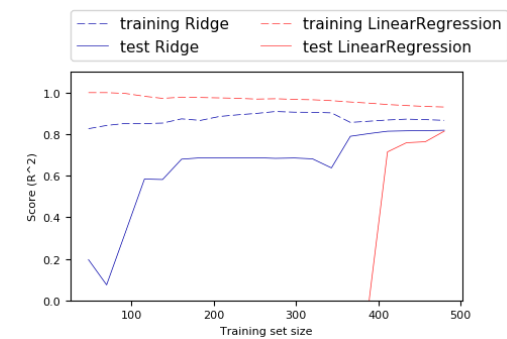
plt.plot(lr.coef_, 'o', label="LinearRegression")
plt.xlabel("Coefficient index")
plt.ylabel("Coefficient magnitude")
plt.hlines(0, 0, len(lr.coef_))
plt.ylim(-25, 25)
plt.legend()
```

Out[9]: <matplotlib.legend.Legend at 0x11cc1ed30>



Another way to understand the influence of regularization is to fix a value of alpha but vary the amount of training data available. With enough training data, regularization becomes less important: ridge and linear regression will have the same performance.

```
In [10]: mglearn.plots.plot_ridge_n_samples  
()
```



Lasso

- Another form of regularization
- Adds a penalty term to the least squares sum : $\lambda \sum_i |w_i|$
- Prefers coefficients to be exactly zero (L1 regularization).
- Some features are entirely ignored by the model: automatic feature selection.
- Same parameter `alpha` to control the strength of regularization.
- New parameter `max_iter`: the maximum number of iterations
 - Should be higher for small values of `alpha`


```
In [11]: from sklearn.linear_model import Lasso

lasso = Lasso().fit(X_train, y_train)
print("Training set score: {:.2f}".format(lasso
.score(X_train, y_train)))
print("Test set score: {:.2f}".format(lasso.sco
re(X_test, y_test)))
print("Number of features used: {}".format(np.s
um(lasso.coef_ != 0)))
```

Training se
t score: 0.
29
Test set sc
ore: 0.21
Number of f
eatures use
d: 4

```
In [12]: # we increase the default setting of "max_ite
r",
# otherwise the model would warn us that we sho
uld increase max_iter.
lasso001 = Lasso(alpha=0.01, max_iter=100000).f
it(X_train, y_train)
print("Training set score: {:.2f}".format(lasso
001.score(X_train, y_train)))
print("Test set score: {:.2f}".format(lasso001.
score(X_test, y_test)))
print("Number of features used: {}".format(np.s
um(lasso001.coef_ != 0)))
```

Training se
t score: 0.
90
Test set sc
ore: 0.77
Number of f
eatures use
d: 33

```
In [13]: lasso00001 = Lasso(alpha=0.0001, max_iter=10000
0).fit(X_train, y_train)
print("Training set score: {:.2f}".format(lasso
00001.score(X_train, y_train))
print("Test set score: {:.2f}".format(lasso0000
1.score(X_test, y_test))
print("Number of features used: {}".format(np.s
um(lasso00001.coef_ != 0)))
```

Training s
et score:
0.95
Test set s
core: 0.64
Number of
features u
sed: 94

We can again analyse what happens to the weights:

```
In [14]: plt.plot(lasso.coef_, 's', label="Lasso alpha=1")
plt.plot(lasso001.coef_, '^', label="Lasso alpha=0.01"
)
plt.plot(lasso00001.coef_, 'v', label="Lasso alpha=0.0
001")
plt.plot(ridge01.coef_, 'o', label="Ridge alpha=0.1")

plt.legend(ncol=2, loc=(0, 1.05))
plt.ylim(-25, 25)
plt.xlabel("Coefficient index")
plt.ylabel("Coefficient magnitude");
```



Interpreting L1 and L2 loss

- Red ellipses are the contours of the least squares error function
- In blue are the constraints imposed by the L1 (left) and L2 (right) loss functions
- For L1, the likelihood of hitting the objective with the corners is higher
 - Weights of other coefficients are 0, hence sparse representations
- For L2, it could intersect at any point, hence non-zero weights
- From *Elements of Statistical Learning*:




FIGURE 3.11. Estimation picture for the lasso (left) and ridge regression (right). Shown are contours of the error and constraint functions. The solid blue areas are the constraint regions $|\beta_1| + |\beta_2| \leq t$ and $\beta_1^2 + \beta_2^2 \leq t^2$, respectively, while the red ellipses are the contours of the least squares error function.

Linear models for Classification

Aims to find a (hyper)plane that separates the examples of each class.
For binary classification (2 classes), we aim to fit the following function:

$$\hat{y} = w_0 * x_0 + w_1 * x_1 + \dots + w_p * x_p + b > 0$$

When $\hat{y} < 0$, predict class -1, otherwise predict class +1

There are many algorithms for learning linear classification models, differing in:

- Loss function: evaluate how well the linear model fits the training data
- Regularization techniques

Most common techniques:

- Logistic regression:
 - `sklearn.linear_model.LogisticRegression`
- Linear Support Vector Machine:
 - `sklearn.svm.LinearSVC`

Logistic regression

Fits a logistic regression curve/surface to the data



- Logistic regression predicts the target using the logarithm of the class probability:

$$Pr[1|x_1, \dots, x_k] \rightarrow \log\left(\frac{Pr[1|x_1, \dots, x_k]}{1 - Pr[1|x_1, \dots, x_k]}\right)$$

- *Logit transformation* maps $[0,1]$ to $(-\text{Inf}, \text{Inf})$
- Resulting class probability (the green curve in the figure above):

$$Pr[1|x_1, \dots, x_k] = \frac{1}{1 + \exp(-(w_0 * x_0 + w_1 * x_1 + \dots + w_p * x_p))}$$

- Parameters w are found from training data using *maximum likelihood*

Maximum likelihood

- Aim: maximize the probability of the observed training data with respect to the final model parameters
- We can use logarithms of probabilities and maximize conditional *log-likelihood* instead of the product of probabilities


$$\sum_{i=1}^n (1 - y^{(i)}) \log(1 - \text{Pr}[1|x_1^{(i)}, \dots, x_k^{(i)}]) + y^{(i)} \text{Pr}[1|x_1^{(i)}, \dots, x_k^{(i)}]$$

- Class values $y^{(i)}$ are either 0 or 1
- Weights w_i need to be chosen to maximize log-likelihood
 - This can be done using *iterative re-weighted least squares*
 - Other optimization methods can be used as well

Linear Support Vector Machine

Find hyperplane maximizing the *margin* between the classes

Linear SVM: Decision Boundary



Prediction is identical to weighted kNN: find the support vector that is nearest, according to a distance measure (kernel) and a weight for each support vector.

Comparison

```
In [15]: from sklearn.linear_model import LogisticRegression
         from sklearn.svm import LinearSVC

         X, y = mglearn.datasets.make_forge()

         fig, axes = plt.subplots(1, 2, figsize=(10, 3))

         for model, ax in zip([LinearSVC(), LogisticRegression()], axes):
             clf = model.fit(X, y)
             mglearn.plots.plot_2d_separator(clf, X, fill=False, eps=0.5,
                                             ax=ax, alpha=.7)
             mglearn.discrete_scatter(X[:, 0], X[:, 1], y, ax=ax)
             ax.set_title("{}\n".format(clf.__class__.__name__))
             ax.set_xlabel("Feature 0")
             ax.set_ylabel("Feature 1")
         axes[0].legend();
```




Both methods can be regularized:

- L2 regularization by default, L1 also possible
- C parameter: inverse of strength of regularization
 - higher C: less regularization
 - penalty for misclassifying points while keeping w_i close to 0

High C values (less regularization): fewer misclassifications but smaller margins.

```
In [16]: mglearn.plots.plot_linear_svc_regularization()
```



Model selection: Logistic regression

```
In [17]: from sklearn.datasets import load_breast_cancer
cancer = load_breast_cancer()
X_train, X_test, y_train, y_test = train_test_split(
    cancer.data, cancer.target, stratify=cancer
    .target, random_state=42)
logreg = LogisticRegression().fit(X_train, y_train)
print("Training set score: {:.3f}".format(logreg
    .score(X_train, y_train)))
print("Test set score: {:.3f}".format(logreg.sc
    ore(X_test, y_test)))
```

Training s
et score:
0.953
Test set s
core: 0.95
8

```
In [18]: logreg100 = LogisticRegression(C=100).fit(X_train, y_train)
print("Training set score: {:.3f}".format(logreg
    100.score(X_train, y_train)))
print("Test set score: {:.3f}".format(logreg100.
    score(X_test, y_test)))
```

Training s
et score:
0.979
Test set s
core: 0.96
5


```
In [19]: logreg001 = LogisticRegression(C=0.01).fit(X_train, y_train)
print("Training set score: {:.3f}".format(logreg
    001.score(X_train, y_train)))
print("Test set score: {:.3f}".format(logreg001.
    score(X_test, y_test)))
```

Training s
et score:
0.934
Test set s
core: 0.93
0

Effect of C on model parameters:

```
In [20]: plt.plot(logreg.coef_.T, 'o', label="C=1")
plt.plot(logreg100.coef_.T, '^', label="C=100")
plt.plot(logreg001.coef_.T, 'v', label="C=0.001")
plt.xticks(range(cancer.data.shape[1]),
cancer.feature_names, rotation=90)
plt.hlines(0, 0, cancer.data.shape[1])
plt.ylim(-5, 5)
plt.xlabel("Coefficient index")
plt.ylabel("Coefficient magnitude")
plt.legend()
```

```
Out[20]: <matplotlib.legend.
Legend at 0x11e
301eb8>
```



Idem with L1 regularization (`penalty='l1'`):


```
In [21]: for C, marker in zip([0.001, 1, 100],
    ['o', '^', 'v']):
    lr_l1 = LogisticRegression(C=C, p
    enalty="l1").fit(X_train, y_train)
    print("Training accuracy of l1 lo
    greg with C={:.3f}: {:.2f}".format(
        C, lr_l1.score(X_train, y_t
    rain)))
    print("Test accuracy of l1 logreg
    with C={:.3f}: {:.2f}".format(
        C, lr_l1.score(X_test, y_te
    st)))
    plt.plot(lr_l1.coef_.T, marker, l
    abel="C={:.3f}".format(C))

plt.xticks(range(cancer.data.shape[1
]), cancer.feature_names, rotation=90
)
plt.hlines(0, 0, cancer.data.shape[1
])
plt.xlabel("Coefficient index")
plt.ylabel("Coefficient magnitude")

plt.ylim(-5, 5)
plt.legend(loc=3)
```

```
Training accuracy of
l1 logreg with C=0.00
1: 0.91
Test accuracy of l1 l
ogreg with C=0.001:
0.92
Training accuracy of
l1 logreg with C=1.00
0: 0.96
Test accuracy of l1 l
ogreg with C=1.000:
0.96
Training accuracy of
l1 logreg with C=100.
000: 0.99
Test accuracy of l1 l
ogreg with C=100.000:
0.98
```

```
Out[21]: <matplotlib.legend.Legend at 0x107f5f400>
```



Linear Models for multiclass classification

Common technique: one-vs.-rest approach:

- A binary model is learned for each class vs. all other classes
- Creates as many binary models as there are classes
- Every binary classifiers makes a prediction, the one with the highest score (>0) wins


Build binary linear models:

```
In [22]: from sklearn.datasets import make_blobs

X, y = make_blobs(random_state=42)
linear_svm = LinearSVC().fit(X, y)

mglearn.discrete_scatter(X[:, 0], X[:, 1], y)
line = np.linspace(-15, 15)
for coef, intercept, color in zip(linear_svm.coef_, linear_svm.intercept_,
                                     mglearn.cm3.colors):
    plt.plot(line, -(line * coef[0] + intercept) / coef[1], c=color)
plt.ylim(-10, 15)
plt.xlim(-10, 8)
plt.xlabel("Feature 0")
plt.ylabel("Feature 1")
plt.legend(['Class 0', 'Class 1', 'Class 2', 'Line class 0', 'Line class 1',
            'Line class 2'], loc=(1.01, 0.3))
```


Out[22]: <matplotlib.legend.
d.Legend at 0x11c
babd30>



Actual predictions (decision boundaries):

```
In [23]: mglearn.plots.plot_2d_classification(linear_svm, X, fill=True, alpha=.7)
mglearn.discrete_scatter(X[:, 0], X[:, 1], y)
line = np.linspace(-15, 15)
for coef, intercept, color in zip(linear_svm.coef_, linear_svm.intercept_,
                                     mglearn.cm3.colors):
    plt.plot(line, -(line * coef[0] + intercept) / coef[1], c=color)
plt.legend(['Class 0', 'Class 1', 'Class 2', 'Line class 0', 'Line class 1',
            'Line class 2'], loc=(1.01, 0.3))
plt.xlabel("Feature 0")
plt.ylabel("Feature 1")
```

Out[23]: <matplotlib.text.Text at 0x11c9e978>



Strengths, weaknesses and parameters

Regularization parameters:

- Regression: alpha (higher values, simpler models)
 - Ridge (L2), Lasso (L1), LinearRegression (None)
- Classification: C (smaller values, simpler models)
 - LogisticRegression or SVC (both have L1/L2 option)

L1 vs L2:

- L2 is default
- Use L1 if you assume that few features are important
 - Or, if model interpretability is important

Other options:

- ElasticNet regression: allows L1 vs L2 trade-off
- SGDClassifier/SGDRegressor: optimize w_i, b with stochastic gradient descent (more scalable)

Consider linear models when:


- number of features is large compared to the number of samples
 - other algorithms perform better in low-dimensional spaces
- very large datasets (fast to train and predict)
 - other algorithms become (too) slow

Intuition: why linear models are powerful in high dimension

While linear models are limited on low-dimensional data, they can often fit high dimensional data very well.

```
In [24]: from sklearn.svm import LinearSVC
X, y = make_blobs(centers=4, random_state=8)
y = y % 2 # Reduces 4 classes to 2
linear_svm = LinearSVC().fit(X, y)

mglearn.plots.plot_2d_separator(linear_svm,
                                X)
mglearn.discrete_scatter(X[:, 0], X[:, 1],
                          y)
plt.xlabel("Feature 0")
plt.ylabel("Feature 1");
```



While in the previous picture the classes (blue and red) cannot be linearly separated, imagine that we have another dimension that tells us more about each class.

```
In [25]: # add the square of the first feature (this is just one way  
         to add new features)  
X_new = np.hstack([X, X[:, 1:] ** 2])  
  
from mpl_toolkits.mplot3d import Axes3D, axes3d  
figure = plt.figure()  
# visualize in 3D  
ax = Axes3D(figure, elev=-152, azimuth=-26)  
# plot first all the points with y==0, then all with y == 1  
mask = y == 0  
ax.scatter(X_new[mask, 0], X_new[mask, 1], X_new[mask, 2], c  
='b',  
           cmap=mglearn.cm2, s=60)  
ax.scatter(X_new[~mask, 0], X_new[~mask, 1], X_new[~mask, 2]  
], c='r', marker='^',  
           cmap=mglearn.cm2, s=60)  
ax.set_xlabel("feature0")  
ax.set_ylabel("feature1")  
ax.set_zlabel("feature1 ** 2");
```


Now we can fit a linear model

Note: We will come back to this when discussing *kernelization*, in which we construct new dimensions on purpose.

```

In [26]: linear_svm_3d = LinearSVC().fit(X_new, y)
coef, intercept = linear_svm_3d.coef_.ravel(
    ), linear_svm_3d.intercept_

# show linear decision boundary
figure = plt.figure()
ax = Axes3D(figure, elev=-152, azimuth=-26)
xx = np.linspace(X_new[:, 0].min() - 2, X_new[:, 0].max() + 2, 50)
yy = np.linspace(X_new[:, 1].min() - 2, X_new[:, 1].max() + 2, 50)

XX, YY = np.meshgrid(xx, yy)
ZZ = (coef[0] * XX + coef[1] * YY + intercept) / -coef[2]
ax.plot_surface(XX, YY, ZZ, rstride=8, cstride=8, alpha=0.3)
ax.scatter(X_new[mask, 0], X_new[mask, 1], X_new[mask, 2], c='b',
           cmap=mglearn.cm2, s=60)
ax.scatter(X_new[~mask, 0], X_new[~mask, 1], X_new[~mask, 2], c='r', marker='^',
           cmap=mglearn.cm2, s=60)


ax.set_xlabel("feature0")
ax.set_ylabel("feature1")
ax.set_zlabel("feature1 ** 2")

```

```

Out[26]: <matplotlib.text.Text at 0x11c8db6a0>

```



Uncertainty estimates from classifiers

Classifiers can often provide uncertainty estimates of predictions.

In practice, you are often interested in how certain a classifier is about each class prediction (e.g. cancer treatments).

Scikit-learn offers 2 functions. Often, both are available for every learner, but not always.

- `decision_function`: returns floating point value for each sample
- `predict_proba`: return probability for each class

In [27]:

```
# create and split a synthetic dataset
from sklearn.linear_model import LogisticRegression
from sklearn.datasets import make_blobs
X, y = make_blobs(centers=2, cluster_std=2.5, random_state=8)

# we rename the classes "blue" and "red"
y_named = np.array(["blue", "red"])[y]

# we can call train test split with arbitrary many arrays
# all will be split in a consistent manner
X_train, X_test, y_train_named, y_test_named, y_train, y_test = \
    train_test_split(X, y_named, y, random_state=0)

# build the logistic regression model
lr = LogisticRegression()
lr.fit(X_train, y_train_named)

mglearn.plots.plot_2d_separator(lr, X)
mglearn.discrete_scatter(X[:, 0], X[:, 1], y);
```



The Decision Function

In the binary classification case, the return value of `decision_function` is of shape `(n_samples,)`, and it returns one floating-point number for each sample. The first class (class 0) is considered negative, the other (class 1) positive.

This value encodes how strongly the model believes a data point to belong to the “positive” class.

- Positive values indicate a preference for the "positive" class
- Negative values indicate a preference for the "negative" (other) class

```
In [28]: # show the first few entries of decision function
print("Decision function:\n{}".format(lr.decision_function(X_test)[:6]))
```

```
Decision function:
[ 0.527  4.314  5.
 92    2.899  4.751
-7.035]
```

```
In [29]: # Recover the predictions by looking at the sign
print("Thresholded decision function:
\n{}".format(
    lr.decision_function(X_test)[:6]
    > 0))
print("Predictions:\n{}".format(lr.predict(X_test)[:6]))
```

```
Thresholded decision
function:
[ True  True  True
 True  True False]
Predictions:
['red' 'red' 'red'
 'red' 'red' 'blue']
```


The range of `decision_function` can be arbitrary, and depends on the data and the model parameters. This makes it sometimes hard to interpret.

```
In [30]: decision_function = lr.decision_function(X_test)
print("Decision function minimum: {:.2f} maximum: {:.2f}".format(
    np.min(decision_function), np.max(decision_function)))
```

```
Decision function minimum: -10.48 maximum: 8.61
```


We can visualize the decision function as follows, with the actual decision boundary left and the values of the decision boundaries color-coded on the right. Note how the test examples are labeled depending on the decision function.

```

In [31]: fig, axes = plt.subplots(1, 2, figsize=(13, 5))

mglearn.tools.plot_2d_separator(lr, X, ax=axes[0], alpha
=.4,
                                fill=True, cm=mglearn.cm
2)
scores_image = mglearn.tools.plot_2d_scores(lr, X, ax=axes[1],
                                alpha=.4, cm
                                =mglearn.ReBl)

for ax in axes:
    # plot training and test points
    mglearn.discrete_scatter(X_test[:, 0], X_test[:, 1],
y_test,
                                markers='^', ax=ax)
    mglearn.discrete_scatter(X_train[:, 0], X_train[:, 1],
y_train,
                                markers='o', ax=ax)
    ax.set_xlabel("Feature 0")
    ax.set_ylabel("Feature 1")
cbar = plt.colorbar(scores_image, ax=axes.tolist())
cbar.set_alpha(1)
cbar.draw_all()
axes[0].legend(["Test class 0", "Test class 1", "Train c
lass 0",
                "Train class 1"], ncol=4, loc=(.1, 1.1
));

```

Predicting probabilities

The output of `predict_proba` is a *probability* for each class, with one column per class. They sum up to 1.

```
In [32]: print("Shape of probabilities: {}".format(lr
            .predict_proba(X_test).shape))
          # show the first few entries of predict_proba
          print("Predicted probabilities:\n{}".format(
                lr.predict_proba(X_test[:6])))
```

```
Shape of probabilities: (2
5, 2)
Predicted probabilities:
[[0.371 0.62
9]
 [0.013 0.98
7]
 [0.003 0.99
7]
 [0.052 0.94
8]
 [0.009 0.99
1]
 [0.999 0.00
1]]
```


We can visualize them again. Note that the gradient looks different now.

```
In [33]: fig, axes = plt.subplots(1, 2, figsize=(13, 5))

mglearn.tools.plot_2d_separator(
    lr, X, ax=axes[0], alpha=.4, fill=True, cm=mglearn.cm2
)
scores_image = mglearn.tools.plot_2d_scores(
    lr, X, ax=axes[1], alpha=.5, cm=mglearn.ReBl, function
    ='predict_proba')

for ax in axes:
    # plot training and test points
    mglearn.discrete_scatter(X_test[:, 0], X_test[:, 1], y
    _test,
                            markers='^', ax=ax)
    mglearn.discrete_scatter(X_train[:, 0], X_train[:, 1],
    y_train,
                            markers='o', ax=ax)
    ax.set_xlabel("Feature 0")
    ax.set_ylabel("Feature 1")
# don't want a transparent colorbar
    cbar = plt.colorbar(scores_image, ax=axes.tolist())
    cbar.set_alpha(1)
    cbar.draw_all()
    axes[0].legend(["Test class 0", "Test class 1", "Train cla
    ss 0",
                  "Train class 1"], ncol=4, loc=(.1, 1.1));
```


Interpreting probabilities

- The class with the highest probability is predicted.
- How well the uncertainty actually reflects uncertainty in the data depends on the model and the parameters.
 - An overfitted model tends to make more certain predictions, even if they might be wrong.
 - A model with less complexity usually has more uncertainty in its predictions.
- A model is called *calibrated* if the reported uncertainty actually matches how correct it is — A prediction made with 70% certainty would be correct 70% of the time.
 - LogisticRegression returns well calibrated predictions by default as it directly optimizes log-loss
 - Linear SVM are not well calibrated. They are *biased* towards points close to the decision boundary.
- Techniques exist (<http://scikit-learn.org/stable/modules/calibration.html>) to calibrate models in post-processing. More in the next lecture.

Compare logistic regression and linear SVM

```

In [34]: from sklearn.svm import SVC
svc = SVC(kernel="linear",C=0.1,probability=True).fit(X_train,
            y_train_named)

fig, axes = plt.subplots(1, 2, figsize=(13, 5))

lr_image = mglearn.tools.plot_2d_scores(
        lr, X, ax=axes[0], alpha=.5, cm=mglearn.ReBl, function
    = 'predict_proba')
svc_image = mglearn.tools.plot_2d_scores(
        svc, X, ax=axes[1], alpha=.5, cm=mglearn.ReBl, function
    = 'predict_proba')

for ax in axes:
    # plot training and test points
    mglearn.discrete_scatter(X_test[:, 0], X_test[:, 1], y
    _test,
                            markers='^', ax=ax)
    mglearn.discrete_scatter(X_train[:, 0], X_train[:, 1],
    y_train,
                            markers='o', ax=ax)
    ax.set_xlabel("Feature 0")
    ax.set_ylabel("Feature 1")
# don't want a transparent colorbar
cbar = plt.colorbar(lr_image, ax=axes.tolist())
cbar.set_alpha(1)
cbar.draw_all()
axes[0].legend(["Test class 0", "Test class 1", "Train class 0",
               "Train class 1"], ncol=4, loc=(.1, 1.1));

```

Uncertainty in multi-class classification

- `decision_function` and `predict_proba` methods also work in the multiclass setting
- always have shape $(n_samples, n_classes)$, except for `decision_function` in the binary case (which only returns the values for the positive class)

Example on the Iris dataset, which has 3 classes

```
In [35]: from sklearn.datasets import load_iris

iris = load_iris()
X_train, X_test, y_train, y_test = train_test_split(
    iris.data, iris.target, random_state=42)

lr2 = LogisticRegression()
lr2 = lr2.fit(X_train, y_train)

print("Decision function:\n{}".format(lr2.decision_function(X_test)[:6, :]))
# show the first few entries of predict_proba
print("Predicted probabilities:\n{}".format(lr2.predict_proba(X_test)[:6]))
```

```
Decision function:
[[ -4.744
 0.102 -1.0
84]
 [ 3.699
-1.937 -10.
976]
 [-10.128
 0.898 4.2
62]
 [ -4.504
-0.5 -0.
92 ]
 [ -4.881
 0.249 -1.5
12]
 [ 3.369
-1.644 -10.
167]]
Predicted probabilities:
[[0.011 0.6
68 0.321]
 [0.886 0.1
14 0. ]
 [0. 0.4
19 0.581]
 [0.016 0.5
61 0.423]
 [0.01 0.7
49 0.241]
```