

# Lecture 2: Linear models

Joaquin Vanschoren, Eindhoven University of Technology


# The Mathematics behind Machine Learning

- To understand machine learning algorithms, it often helps to describe them mathematically.
- To avoid confusion, let's specify a precise notation

## Basic notation

- A *scalar* is a simple numeric value, denoted by italic letter:  $x = 3.24$
- A *vector* is a 1D ordered array of  $n$  scalars, denoted by bold letter:  
 $\mathbf{x} = [3.24, 1.2]$ 
  - A vector can represent a *point* in an  $n$ -dimensional space, given a *basis*.
  - $x_i$  denotes the  $i$ th element of a vector, thus  $x_0 = 3.24$ .
    - Note: some other courses use  $x^{(i)}$  notation
- A *set* is an *unordered* collection of unique elements, denote by caligraphic capital:  $S = \{3.24, 1.2\}$
- A *matrix* is a 2D array of scalars, denoted by bold capital:  
 $\mathbf{X} = \begin{bmatrix} 3.24 & 1.2 \\ 2.24 & 0.2 \end{bmatrix}$ 
  - It can represent a set of points in an  $n$ -dimensional space, given a *basis*.
  - $\mathbf{X}_i$  denotes the  $i$ th *row* of the matrix
  - $\mathbf{X}_{i,j}$  denotes the *element* in the  $i$ th row,  $j$ th column, thus  
 $\mathbf{X}_{0,1} = 2.24$
- The *standard basis* for a Euclidean space is the set of unit vectors
  - Data can also be represented in a non-standard basis (e.g. polynomials) if useful

- A *tensor* is an  $k$ -dimensional array of data, denoted by an italic capital:  $T$ 
  - $k$  is also called the *order*, *degree*, or *rank*
  - $T_{i,j,k,\dots}$  denotes the element or sub-tensor in the corresponding position
  - A set of color images can be represented by:
    - a 4D tensor (sample x height x weight x color channel)
    - a 2D tensor (sample x flattened vector of pixel values)



# Basic operations

- Sums and products are denoted by capital Sigma and capital Pi:

$$\sum_{i=0}^n = x_0 + x_1 + \dots + x_p \quad \prod_{i=0}^n = x_0 \cdot x_1 \cdot \dots \cdot x_p$$

- Operations on vectors are *element-wise*: e.g.

$$\mathbf{x} + \mathbf{z} = [x_0 + z_0, x_1 + z_1, \dots, x_p + z_p]$$

- Dot product

$$\mathbf{w}\mathbf{x} = \mathbf{w} \cdot \mathbf{x} = \sum_{i=0}^p w_i \cdot x_i = w_0 \cdot x_0 + w_1 \cdot x_1 + \dots + w_p \cdot x_p$$

- Matrix product  $\mathbf{W}\mathbf{x} = \begin{bmatrix} \mathbf{w}_0 \cdot \mathbf{x} \\ \dots \\ \mathbf{w}_p \cdot \mathbf{x} \end{bmatrix}$

- A function  $f(x) = y$  relates an input element  $x$  to an output  $y$

- It has a *local minimum* at  $x = c$  if  $f(x) \geq f(c)$  in interval  $(c - \epsilon, c + \epsilon)$

- It has a *global minimum* at  $x = c$  if  $f(x) \geq f(c)$  for any value for  $x$

- A vector function consumes an input and produces a vector:  $\mathbf{f}(\mathbf{x}) = \mathbf{y}$

- $\max_{x \in X} f(x)$  returns the highest value  $f(x)$  for any  $x$


- $\operatorname{argmax}_{c \in C} f(x)$  returns the element  $c$  that maximizes  $f(c)$

# Gradients

- A *derivative*  $f'$  of a function  $f$  describes how fast  $f$  grows or decreases
- The process of finding a derivative is called differentiation
  - Derivatives for basic functions are known
  - For non-basic functions we use the *chain rule*:
$$F(x) = f(g(x)) \rightarrow F'(x) = f'(g(x))g'(x)$$
- A function is *differentiable* if it has a derivate in any point of it's domain
  - It's *continuously differentiable* if  $f'$  is itself a function
  - It's *smooth* if  $f', f'', f''', \dots$  all exist
- A *gradient*  $\nabla f$  is the derivate of a function in multiple dimensions
  - It is a vector of *partial derivatives*:  $\nabla f = \left[ \frac{\partial f}{\partial x_0}, \frac{\partial f}{\partial x_1}, \dots \right]$
  - E.g.  $f = 2x_0 + 3x_1^2 - \sin(x_2) \rightarrow \nabla f = [2, 6x_1, -\cos(x_2)]$

# Probabilities

- A random variable  $X$  can be continuous or discrete
- A probability distribution of a discrete variable is the list of probabilities for each possible value
  - Also called the *probability mass function* (pmf)
  - The *expectation* (or *mean*)  $\mathbb{E}[X] = \mu_X = \sum_{i=1}^k [x_i \cdot \Pr(X = x_i)]$
  - The standard deviation  $\sigma = \sqrt{\mathbb{E}[(X - \mu)^2]}$  and variance  $\text{var}(X) = \sigma^2$
- A probability distribution  $f_X$  of a continuous variable  $X$  is described by a *probability density function* (pdf)
  - The *expectation* is given by  $\mathbb{E}[X] = \int_{\mathbb{R}} x f_X(x) dx$



# Linear models

Linear models make a prediction using a linear function of the input features.  
Can be very powerful for or datasets with many features.

If you have more features than training data points, any target  $y$  can be perfectly modeled (on the training set) as a linear function.




# Linear models for regression

Prediction formula for input features  $\mathbf{x}$ .  $w_i$  and  $b$  are the *model parameters* that need to be learned.

$$\hat{y} = \mathbf{w}\mathbf{x} + b = \sum_{i=0}^p w_i \cdot x_i + b = w_0 \cdot x_0 + w_1 \cdot x_1 + \dots + w_p \cdot x_p + b$$

There are many different algorithms, differing in how  $w$  and  $b$  are learned from the training data.

w[0]: 0.393906    b: -0.031804




# Linear Regression aka Ordinary Least Squares

- Finds the parameters  $w$  and  $b$  that minimize the *mean squared error* between predictions (red) and the true regression targets (blue),  $y$ , on the training set.
  - MSE: Sum of the squared differences (residuals) between the predictions  $\hat{y}_i$  and the true values  $y_i$ .

$$\mathcal{L}_{MSE} = \sum_{n=0}^N (y_n - \hat{y}_n)^2 = \sum_{n=0}^N (y_n - (\mathbf{w}\mathbf{x}_n + b))^2$$

- Convex optimization problem with unique closed-form solution (if you have more data points than model parameters  $w$ )
- It has no hyperparameters, thus model complexity cannot be controlled.



Linear regression can be found in `sklearn.linear_model`. We'll evaluate it on the Boston Housing dataset.

```
lr = LinearRegression().fit(X_train, y_train)
```

```
Weights (coefficients): [ -412.711   -52.243  -131.899   -12.004   -15.511
 28.716    54.704
   -49.535    26.582    37.062   -11.828   -18.058   -19.525    12.203
 2980.781  1500.843   114.187   -16.97    40.961   -24.264    57.616
 1278.121 -2239.869   222.825    -2.182   42.996   -13.398   -19.389
    -2.575   -81.013     9.66     4.914    -0.812    -7.647    33.784
   -11.446    68.508   -17.375   42.813     1.14 ]
Bias (intercept): 30.93456367364464
```

```
Training set score (R^2): 0.95
```

```
Test set score (R^2): 0.61
```

# Ridge regression

- Same formula as linear regression
- Adds a penalty term to the least squares sum:

$$\mathcal{L}_{Ridge} = \sum_{n=0}^N (y_n - (\mathbf{w}\mathbf{x}_n + b))^2 + \alpha \sum_{i=0}^p w_i^2$$

- Requires that the coefficients (w) are close to zero.
  - Each feature should have as little effect on the outcome as possible
- Regularization: explicitly restrict a model to avoid overfitting.
- Type of L2 regularization: prefers many small weights
  - L1 regularization prefers sparsity: many weights to be 0, others large

```
Ridge can also be found in sklearn.linear_model.  
ridge = Ridge().fit(X_train, y_train)
```

```
Training set score: 0.89  
Test set score: 0.75
```



Test set score is higher and training set score lower: less overfitting!

The strength of the regularization can be controlled with the `alpha` parameter. Default is 1.0.


- Increasing `alpha` forces coefficients to move more toward zero (more regularization)
- Decreasing `alpha` allows the coefficients to be less restricted (less regularization)

We can plot the weight values for different levels of regularization. Move the slider to increase/decrease regularization. Increasing regularization decreases the values of the coefficients, but never to 0.





Another way to understand the influence of regularization is to fix a value of alpha but vary the amount of training data available. With enough training data, regularization becomes less important: ridge and linear regression will have the same performance.



# Lasso (Least Absolute Shrinkage and Selection Operator)

- Another form of regularization
- Adds a penalty term to the least squares sum:

$$\mathcal{L}_{Lasso} = \sum_{n=0}^N (y_n - (\mathbf{w}\mathbf{x}_n + b))^2 + \alpha \sum_{i=0}^p |w_i|$$


- Prefers coefficients to be exactly zero (L1 regularization).
- Some features are entirely ignored by the model: automatic feature selection.
- Same parameter `alpha` to control the strength of regularization.
- Convex, but no longer strictly convex (and NOT differentiable). Weights can be optimized using (for instance) *coordinate descent*
- New parameter `max_iter`: the maximum number of coordinate descent iterations
  - Should be higher for small values of `alpha`

# Gradient Descent

- Start with a random set of  $p$  weights values  $\mathbf{w}^0$
- Compute the derivative of the objective function  $\mathcal{L}$  (e.g.  $\mathcal{L}_{\mathcal{R}idge}$ ) and use it to find the slope (in  $p$  dimensions)
- Update all weights slightly (step size  $\gamma$ ) in the direction of the downhill slope. For step  $s$ :

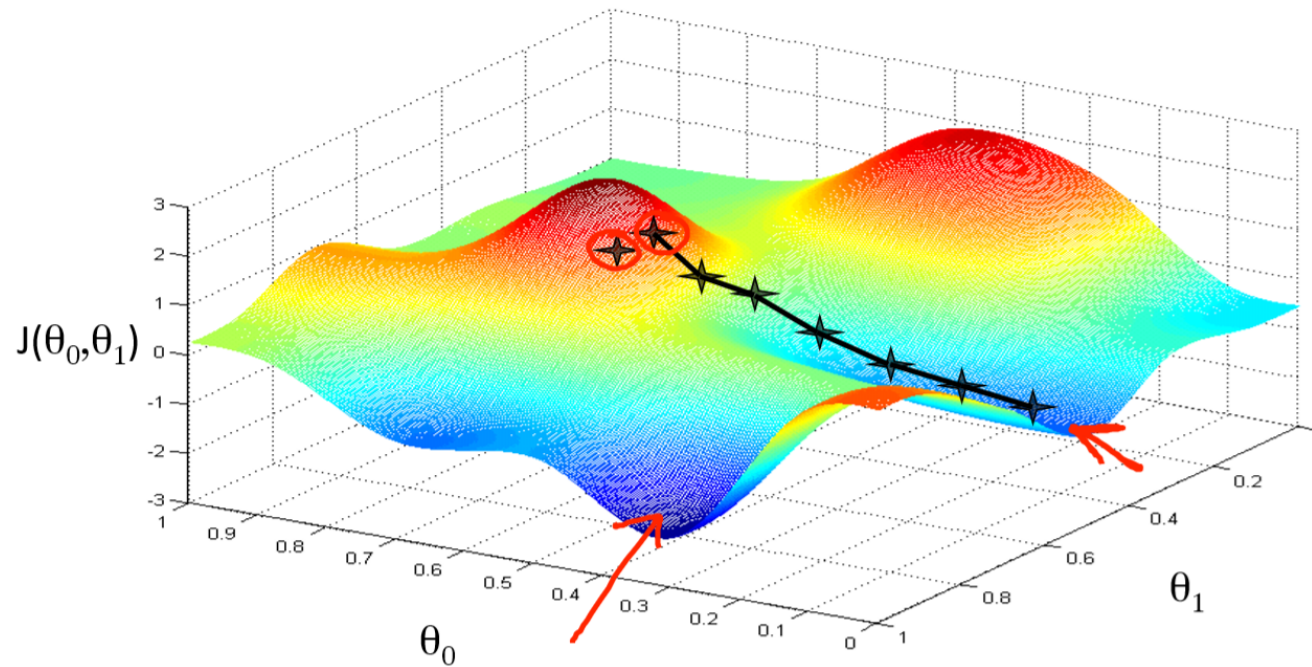
$$\mathbf{w}^{s+1} = \mathbf{w}^s - \gamma \nabla \mathcal{L}(\mathbf{w}^s)$$

- Repeat for `max_iter` iterations
- Visualization in 1 dimension (for 1 weight):




# Gradient Descent

In two dimensions:




# Gradient Descent

- Intuition: walking downhill using only the slope you "feel" nearby





# Coordinate descent

- Variation of gradient descent, also applicable for non-differentiable loss functions
- Faster iterations, may converge more slowly
- In every iteration, optimizes a single coordinate  $w_i$ , using a coordinate selection rule (e.g. round robin)
  - If  $\mathcal{L}$  is differentiable, the update rule is based on the partial derivative:  $w_i^{s+1} = w_i^s - \gamma \frac{\partial \mathcal{L}}{\partial w_i}$
  - If  $\mathcal{L}$  is not differentiable but convex, the subgradient (<https://www.cs.cmu.edu/~ggordon/10725-F12/slides/06-sg-method.pdf>) can be computed.
- For Lasso, the resulting update rule ([https://xavierbourretsicotte.github.io/lasso\\_derivation.html](https://xavierbourretsicotte.github.io/lasso_derivation.html)) includes the *soft thresholding operator*  $S$ :  $w_i^{s+1} \cong S(f(w_i^s), \alpha)$ 
  - $S$  sets  $w_i$ 's to 0 when they are sufficiently small ('sufficiently' is defined by  $\alpha$ )






We can again analyse what happens to the weights. Increasing regularization under L1 leads to many coefficients becoming exactly 0.



# Interpreting L1 and L2 loss

- Red ellipses are the contours of the least squares error function
- In blue are the constraints imposed by the L1 (left) and L2 (right) loss
- For L1, the likelihood of hitting the objective with the corners is higher
  - Weights of other coefficients are 0, hence sparse representations
- For L2, it could intersect at any point, hence non-zero weights
- From *Elements of Statistical Learning*:



**FIGURE 3.11.** Estimation picture for the lasso (left) and ridge regression (right). Shown are contours of the error and constraint functions. The solid blue areas are the constraint regions  $|\beta_1| + |\beta_2| \leq t$  and  $\beta_1^2 + \beta_2^2 \leq t^2$ , respectively, while the red ellipses are the contours of the least squares error function.

## Linear models for Classification

Aims to find a (hyper)plane that separates the examples of each class.  
For binary classification (2 classes), we aim to fit the following function:

$$\hat{y} = w_0 * x_0 + w_1 * x_1 + \dots + w_p * x_p + b > 0$$

When  $\hat{y} < 0$ , predict class -1, otherwise predict class +1

There are many algorithms for learning linear classification models, differing in:

- Loss function: evaluate how well the linear model fits the training data
- Regularization techniques

Most common techniques:

- Logistic regression:
  - `sklearn.linear_model.LogisticRegression`
- Linear Support Vector Machine:
  - `sklearn.svm.LinearSVC`


## *Logistic regression*

The logistic model uses the *logistic* (or *sigmoid*) function to estimate the probability that a given sample belongs to class 1:

$$z = f(x) = w_0 * x_0 + w_1 * x_1 + \dots + w_p * x_p$$
$$\hat{y} = Pr[1|x_1, \dots, x_k] = g(z) = \frac{1}{1 + e^{-z}}$$



On 2-dimensional data:



- The logistic function is chosen because it maps values  $(-\infty, \infty)$  to a probability  $[0, 1]$
- We add a new dimension for the dependent variable  $y$  and fit the logistic function  $g(z)$  so that it separates the samples as good as possible. The positive (blue) points are mapped to 1 and the negative (red) points to 0.
- After fitting, the logistic function provides the probability that a new point is positive. If we need a binary prediction, we can threshold at 0.5.
- There are different ways to find the optimal parameters  $w$  that fit the training data best



### *Fitting (solving): cross-entropy*

- We define the difference (error) between the actual probabilities (frequencies)  $p_i$  and the predicted probabilities  $q_i$  is the cross-entropy  $H(p, q)$ :

$$H(p, q) = - \sum_i p_i \log(q_i)$$

- Note: Instead of minimizing cross-entropy  $H(p, q)$ , you can maximize *log-likelihood*  $-H(p, q)$ , and hence this is also called *maximum likelihood estimation*
- In binary classification,  $i = 0, 1$  and  $p_1 = y, p_0 = 1 - y, q_1 = \hat{y}, q_0 = 1 - \hat{y}$
- And thus:

$$H(p, q) = -y \log(\hat{y}) - (1 - y) \log(1 - \hat{y})$$

### *Fitting (solving): cross-entropy loss*

- Loss function: the average of all cross-entropies in the sample (of  $N$  data points):

$$\mathcal{L}_{log}(\mathbf{w}) = \sum_{n=1}^N H(p_n, q_n) = \sum_{n=1}^N \left[ -y_n \log(\hat{y}_n) - (1 - y_n) \log(1 - \hat{y}_n) \right]$$

with

$$\hat{y}_n = \frac{1}{1 + e^{\mathbf{w} \cdot \mathbf{x}}}$$

- This is called *logistic loss*, *log loss* or *cross-entropy loss*
- We can (and should always) add a regularization term, either L1 or L2, e.g. for L2:

$$\mathcal{L}_{log}'(\mathbf{w}) = \mathcal{L}_{log}(\mathbf{w}) + \alpha \sum_i w_i^2$$


- Note: sklearn uses  $C$  instead of  $\alpha$ , and it is the inverse (smaller values, more regularization)

## *Fitting (solving): optimization methods*

- There are different ways to optimize cross-entropy loss.
- Gradient descent
  - The logistic function is differentiable, so we can use (stochastic) gradient descent
  - Stochastic Average Gradient descent (SAG): only updates gradient in one direction at each step
- Coordinate descent (default, called `liblinear` in sklearn)
  - Faster, may converge more slowly, may more easily get stuck in local minima
- Newton-Rhapson (or Newton Conjugate Gradient):
  - Finds optima by computing second derivatives (more expensive)
  - Works well if solution space is (near) convex
  - Also known as *iterative re-weighted least squares*
- Quasi-Newton methods
  - Approximate, faster to compute
  - E.g. Limited-memory Broyden–Fletcher–Goldfarb–Shanno (`lbfgs`)

## *Linear Support Vector Machine (intuition)*


Find hyperplanes (dashed lines) maximizing the *margin* between the classes



## *Optimization and prediction*

- Prediction is identical to (weighted) kNN:
  - Points closest to the red support vector are classified red, others blue
  - A support vector can also have a weight (see later)
- The objective function penalizes every point predicted to be on the wrong side of its hyperplane
  - This is called *hinge loss*
- This results in a convex optimization problem solved using the *Langrange Multipliers* method
  - Can also be solved using gradient descent
- This will all be discussed at length later


## Comparison



Both methods can be regularized:

- L2 regularization by default, L1 also possible
- C parameter: inverse of strength of regularization
  - higher C: less regularization
  - penalty for misclassifying points while keeping  $w_i$  close to 0

High  $C$  values (less regularization): fewer misclassifications but smaller margins.







Model selection: Logistic regression

```
logreg = LogisticRegression(C=1).fit(X_train, y_train)
```

Adjust the slider to see the effect of C and L1/L2 regularization



# Uncertainty estimates from classifiers

- Classifiers can often provide uncertainty estimates of predictions.
- Remember that linear models actually return a numeric value.
  - When  $\hat{y} < 0$ , predict class -1, otherwise predict class +1
- In practice, you are often interested in how certain a classifier is about each class prediction (e.g. cancer treatments).

$$\hat{y} = w_0 * x_0 + w_1 * x_1 + \dots + w_p * x_p + b$$


Scikit-learn offers 2 functions. Often, both are available for every learner, but not always.

- `decision_function`: returns floating point value for each sample
- `predict_proba`: return probability for each class


# The Decision Function

In the binary classification case, the return value of `decision_function` is of shape `(n_samples,)`, and it returns one floating-point number for each sample. This value encodes how strongly the model believes a data point to belong to the "positive" class.

- Positive values indicate a preference for the "positive" class
- Negative values indicate a preference for the "negative" (other) class



- The range of `decision_function` can be arbitrary, and depends on the data and the model parameters. This makes it sometimes hard to interpret.
- We can visualize the decision function as follows, with the actual decision boundary left and the values of the decision boundaries color-coded on the right.
- Note how the test examples are labeled depending on the decision function.



## Predicting probabilities


The output of `predict_proba` is a *probability* for each class, with one column per class. They sum up to 1.

```
Shape of probabilities: (25, 2)
```

```
Predicted probabilities:
```

```
[[0.371 0.629]  
 [0.013 0.987]  
 [0.003 0.997]  
 [0.052 0.948]  
 [0.009 0.991]  
 [0.999 0.001]]
```

We can visualize them again. Note that the gradient looks different now.




# Interpreting probabilities

- The class with the highest probability is predicted.
- How well the uncertainty actually reflects uncertainty in the data depends on the model and the parameters.
  - An overfitted model tends to make more certain predictions, even if they might be wrong.
  - A model with less complexity usually has more uncertainty in its predictions.
- A model is called *calibrated* if the reported uncertainty actually matches how correct it is — A prediction made with 70% certainty would be correct 70% of the time.
  - LogisticRegression returns well calibrated predictions by default as it directly optimizes log-loss
  - Linear SVM are not well calibrated. They are *biased* towards points close to the decision boundary.
- Calibration techniques (<http://scikit-learn.org/stable/modules/calibration.html>) can calibrate models in post-processing.



## Compare logistic regression and linear SVM




# Linear Models for multiclass classification


Common technique: one-vs.-rest approach:

- A binary model is learned for each class vs. all other classes
- Creates as many binary models as there are classes
- Every binary classifiers makes a prediction, the one with the highest score ( $>0$ ) wins

Build binary linear models:



Actual predictions (decision boundaries):



# Uncertainty in multi-class classification

- `decision_function` and `predict_proba` also work in the multiclass setting
- always have shape (n\_samples, n\_classes)
- Example on the Iris dataset, which has 3 classes:

Decision function:

```
[[ -4.748   0.113  -1.084]
 [  3.683  -1.914 -10.976]
 [-10.126   0.901   4.262]
 [ -4.505  -0.498  -0.92 ]
 [ -4.887   0.265  -1.512]
 [  3.354  -1.622 -10.167]]
```

Predicted probabilities:

```
[[0.011 0.669 0.32 ]
 [0.884 0.116 0.    ]
 [0.     0.419 0.581]
 [0.016 0.561 0.423]
 [0.01  0.75  0.24 ]
 [0.854 0.146 0.    ]]
```

# Algorithm overview

| Name                | Representation       | Loss function | Optimization                | Regularization  |
|---------------------|----------------------|---------------|-----------------------------|-----------------|
| Least squares       | Linear function      | MSE           | Convex optimization (CFS)   | None            |
| Ridge regression    | Linear function      | MSE + L2      | Convex optimization (CFS)   | L2 loss (alpha) |
| Lasso regression    | Linear function      | MSE + L1      | Coordinate descent          | L1 loss (alpha) |
| Logistic regression | Linear function      | Log Loss      | Gradient/coordinate descent | L1 or L2 (C)    |
| Linear SVM          | Weighted data points | Hinge loss    | Constrained optimization    | L1 or L2 (C)    |
| kNN                 | All data points      | N/A           | N/A                         | n_neighbors     |

# Linear models: Strengths, weaknesses and parameters

Regularization parameters:

- Regression: alpha (higher values, simpler models)
  - Ridge (L2), Lasso (L1), LinearRegression (None)
- Classification: C (smaller values, simpler models)
  - LogisticRegression or SVC (both have L1/L2 option)

L1 vs L2:

- L2 is default
- Use L1 if you assume that few features are important
  - Or, if model interpretability is important

Other options:

- ElasticNet regression: allows L1 vs L2 trade-off
- SGDClassifier/SGDRegressor: optimize  $w_i, b$  with stochastic gradient descent (more scalable)

Consider linear models when:

- number of features is large compared to the number of samples
  - other algorithms perform better in low-dimensional spaces
- very large datasets (fast to train and predict)
  - other algorithms become (too) slow



# Summary

- Linear models
  - Go-to as a first algorithm to try, good for very large datasets, good for very high-dimensional data.
- Regularization is important. Choose between Ridge (L2) or Lasso (L1)
- Each algorithm has its own objective function, and each function can be optimized by certain techniques.
- Regression:
  - Ridge: L2 loss + least squares
  - Lasso: L1 loss + gradient descent
- Classification:
  - Logistic regression: Cross-entropy + gradient descent (or others)
  - SVM: Hinge loss + Lagrange multipliers
- Classifiers return a certainty estimate. The actual prediction can be *calibrated* towards a specific goal.
- Multi-class classification can be done using a one-vs-all approach