

Lecture 3: Model Selection

Joaquin Vanschoren, Eindhoven University of Technology

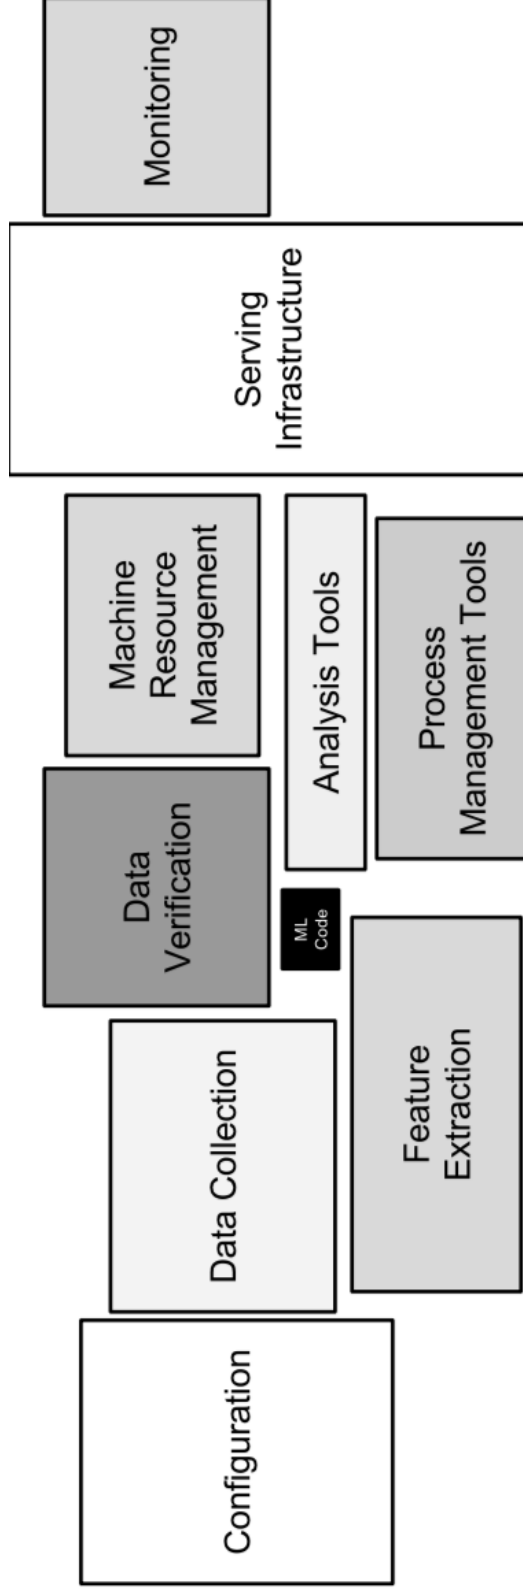
Evaluation

- To know whether we can *trust* our method or system, we need to evaluate it.
- If you cannot measure it, you cannot improve it.
- Model selection: choose between different models in a data-driven way.
- Convince others that your work is meaningful
 - Peers, leadership, clients, yourself(!)
- Keep evaluating relentlessly, adapt to changes

Designing Machine Learning systems

- Just running your favourite algorithm is usually not a great way to start
- Consider the problem at large
 - Do you want to understand phenomena or do black box modelling?
 - How to define and measure success? Are there costs involved?
 - Do you have the right data? How can you make it better?
- Build prototypes early-on to evaluate the above.


- Analyze your model's mistakes
 - Should you collect more, or additional data?
 - Should the task be reformulated?
 - Often a higher payoff than endless finetuning
- Technical debt: creation-maintenance trade-off
 - Very complex machine learning systems are hard/impossible to put into practice
 - See 'Machine Learning: The High Interest Credit Card of Technical Debt'



Only a small fraction of real-world ML systems is composed of the ML code

Real world evaluations

- Evaluate predictions, but also how outcomes improve *because of them*
- Feedback loops: predictions are fed into the inputs, e.g. as new data, invalidating models
- The signal your model found may just be an artifact of your biased data
 - When possible, try to *interpret* what your model has learned
 - See 'Why Should I Trust You?' by Marco Ribeiro et al.



(a) Husky classified as wolf

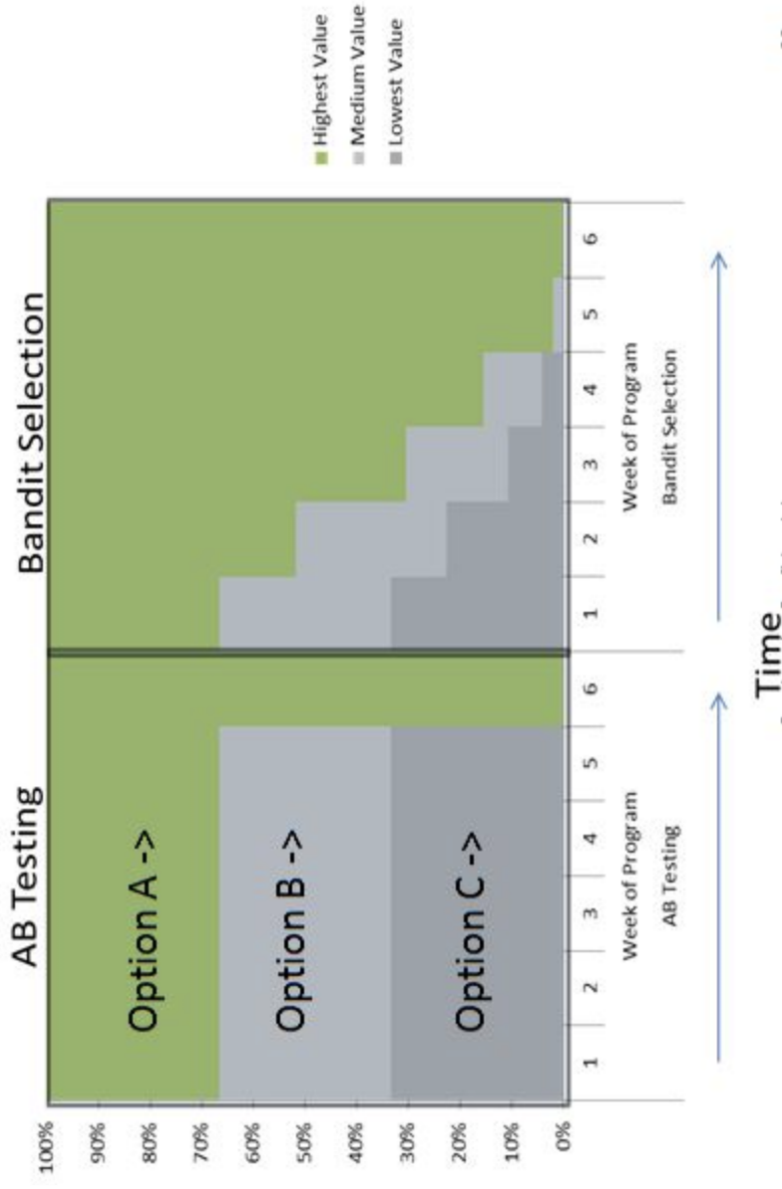


(b) Explanation

- Adversarial situations (e.g. spam filtering) can subvert your predictions
- Do A/B testing (or bandit testing) to evaluate algorithms in the wild

A/B and bandit testing

- Test a single innovation (or choose between two models)
- Have most users use the old system, divert small group to new system
- Evaluate and compare performance
- Bandit testing: smaller time intervals, direct more users to currently winning system



Performance estimation techniques

- We do not have access to future observations
- Evaluate models *as if they are predicting the future*
- Set aside data for objective evaluation
 - How?

The holdout (simple train-test split)


We've already seen the most basic form of evaluation:

- *Randomly* split data (and corresponding labels) into training and test set (75%-25%)
- Train (fit) a model on the training data
- Score a model on the test data (comparing predicted and true labels)
 - We are interested in how well the model *generalizes* to new (test) data

In scikit-learn: train_test_split


```
X, y = make_blobs(centers=2, random_state=0)
# split data and labels into a training and a test set
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)
)
# Fit a model to the training set
model = LogisticRegression().fit(X_train, y_train)
# Evaluate on the test data
test_score = model.score(X_test, y_test)
```

Test set score: 0.92



Cross-validation

- What if one random split yields different models (and scores) than another?
- Reduce bias by testing on every point exactly once
- *k-fold cross-validation* (CV): split (randomized) data into k equal-sized parts, called *folds*
 - First, fold 1 is the test set, and folds 2-5 comprise the training set
 - Then, fold 2 is the test set, folds 1,3,4,5 comprise the training set
 - Compute k evaluation scores, aggregate afterwards (e.g. take the mean)



In scikit-learn:

- `cross_val_score` function with learner, data, labels, number of folds
- Returns list of all scores. Models are built internally, but not returned
- Defaults: 3-fold CV, accuracy (classification) or R^2 (regression)
- Note that there can be quite some *variance* in the results
 - Depends on the stability of the model and the amount of training data
 - Typically, the more training data, the more stable the models

```
logreg = LogisticRegression()  
scores = cross_val_score(logreg, iris.data, iris.target, cv=5)
```


```
Cross-validation scores: [1.  0.967 0.933 0.9  1.  ]  
Average cross-validation score: 0.96  
Variance in cross-validation score: 0.0015
```

Benefits and drawbacks of cross-validation

- More robust: every training example will be in a test set exactly once
 - Model is evaluated on all samples, needs to do well on all
 - With a train-test split, we can be
 - 'lucky': all easy examples in test set
 - 'unlucky': all hard examples in test set
- Shows how *sensitive* the model is to the exact training set
- Better estimation of true performance
 - 10-fold CV uses 90% of all data for training (vs 75% for holdout)
 - The higher k , the more accurate the estimation
- Disadvantage: computational cost, roughly k times slower than holdout
- Unstable models (e.g. deep learning) may not converge for every fold

Stratified K-Fold cross-validation

- If the data is *unbalanced*, some classes have many fewer samples
- Likely that some classes are not present in the test set
- Stratification: *proportions* between classes are conserved in each fold
 - Order examples per class
 - Separate the samples of each class in k sets (strata)
 - Combine corresponding strata into folds



In scikit-learn:

- Uses stratified cross-validation by default for classification
- Normal cross-validation for regression
- Both are non-randomized (samples are not shuffled beforehand)
 - ordered data (e.g. time series) should never be randomized
- You can build folds manually with KFold or StratifiedKFold
 - randomizable (shuffle parameter)

```
kfold = KFold(n_splits=5, shuffle=False) # Not stratified
cross_val_score(logreg, iris.data, iris.target, cv=kfold)
skfold = StratifiedKFold(n_splits=5, shuffle=True)
cross_val_score(logreg, iris.data, iris.target, cv=skfold)
```

```
Cross-validation scores KFold(n_splits=5):
[1.  0.933 0.433 0.967 0.433]
Cross-validation scores StratifiedKFold(n_splits=5, shuffle=True):
[1.  1.  0.867 0.967 1.  ]
```

Can you explain this result?

```
kfold = KFold(n_splits=3)
print("Cross-validation scores KFold(n_splits=3):\n{}".format(
    cross_val_score(logreg, iris.data, iris.target, cv=kfold)))
```

```
Cross-validation scores KFold(n_splits=3):
[0. 0. 0.]
```

Leave-One-Out cross-validation


- k fold cross-validation with k equal to the number of samples
- Completely unbiased (in terms of data splits), but computationally expensive
- But: generalizes *less* well towards unseen data
 - The training sets are correlated (overlap heavily)
 - Overfits on the data used for (the entire) evaluation
 - A different sample of the data can yield different results
- Recommended only for small datasets

```
loo = LeaveOneOut()  
scores = cross_val_score(logreg, iris.data, iris.target, cv=loo)
```

```
Number of cv iterations: 150  
Mean accuracy: 0.95
```


Shuffle-Split cross-validation

- Samples a number of samples (`train_size`) randomly as the training set, and a disjoint number of samples (`test_size`) as the test set
- Repeat this procedure `n_iter` times, obtaining `n_iter` scores
- Handy when using very large datasets
- Example with `train_size = 5`, `test_size = 2`, `n_iter = 4`



In scikit-learn:

- `ShuffleSplit` and `StratifiedShuffleSplit` (recommended for classification)
- `train_size` and `test_size` can be absolute numbers or a percentage of the total dataset

```
shuffle_split = StratifiedShuffleSplit(test_size=.5, train_size=.5, n_splits=10)
scores = cross_val_score(logreg, iris.data, iris.target, cv=shuffle_split)
```

Cross-validation scores:

```
[0.907 0.96 0.92 0.96 0.973 0.96 0.973 0.96 0.92 0.96 ]
```

Note: this is related to *bootstrapping*:

- Sample n (total number of samples) data points, with replacement, as training set (the bootstrap)
- Use the unsampled (out-of-bootstrap) samples as the test set
- Repeat n_iter times, obtaining n_iter scores
- Not supported in scikit-learn, use Shuffle-Split instead.
 - With `train_size=0.66, test_size=0.34`
 - You can prove that bootstraps include 66% of all data points on average

Repeated cross-validation

- Cross-validation is still biased in that the initial split can be made in many ways
- Repeated, or n-times-k-fold cross-validation:
 - Shuffle data randomly, do k-fold cross-validation
 - Repeat n times, yields n times k scores
- Unbiased, very robust, but n times more expensive

Cross-validation with groups

- Sometimes the data contains inherent groups:
 - Blood analysis results on specific patients
 - Facial expressions of specific people
- With normal cross-validation, data from the same person may end up in the training *and* test set
- We want to measure how well the model generalizes to *other* people
- We want to make sure that data points from one person are in *either* the training or test set
 - This is called *grouping* or *blocking*
 - Leave-one-subject-out cross-validation: create test set for each user individually

In scikit-learn

- Add an array with group membership to `cross_val_scores`
 - Use `GroupKFold` with the number of groups as CV procedure
- ```
groups = [0, 0, 0, 1, 1, 1, 2, 2, 2, 3, 3, 3]
scores = cross_val_score(logreg, x, y, groups, cv=GroupKFold(n_splits=4))
```

```
cross_val_score(logreg, x, y, groups, cv=GroupKFold(n_splits=4))
Cross-validation scores :
[0.667 0.667 1. 0.667]
```

# Choosing a performance estimation procedure

No strict rules, only guidelines:

- Always use stratification for classification
- Use holdout for very large datasets (e.g. >1.000.000 examples)
  - Or when learners don't always converge (e.g. deep learning)
- Choose  $k$  depending on dataset size and resources
  - Use leave-one-out for small datasets (e.g. <500 examples)
  - Use cross-validation otherwise
    - Most popular (and theoretically sound): 10-fold CV
    - Literature suggests 5x2-fold CV is better
- Use grouping or leave-one-subject-out for grouped data

# Evaluation Metrics and scoring

Keep the end-goal in mind



# Evaluation vs Optimization

- Each algorithm optimizes a given objective function (on the training data)

- E.g. remember L2 loss in Ridge regression

$$\mathcal{L}_{ridge} = \sum_i (y_i - \sum_j x_{i,j} w_j)^2 + \alpha \sum_i w_i^2$$

- The choice of function is limited by what can be efficiently optimized
  - E.g. gradient descent requires a differentiable loss function
- We *evaluate* the resulting model with a score that makes sense in the real world
  - E.g. percentage of correct predictions (on a test set)
- We also tune the algorithm's hyperparameters to maximize that score

# Binary classification

- We have a positive and a negative class
- 2 different kind of errors:
  - False Positive (type I error): model predicts positive while the true label is negative
  - False Negative (type II error): model predicts negative while the true label is positive
- They are not always equally important
  - Which side do you want to err on for a medical test?

## Confusion matrices

- We can represent all predictions (correct and incorrect) in a confusion matrix
  - n by n array (n is the number of classes)
  - Rows correspond to true classes, columns to predicted classes
  - Each entry counts how often a sample that belongs to the class corresponding to the row was classified as the class corresponding to the column.
  - For binary classification, we label these true negative (TN), true positive (TP), false negative (FN), false positive (FP)

|                |                    |                    |
|----------------|--------------------|--------------------|
| negative class | TN                 | FP                 |
| positive class | FN                 | TP                 |
|                | predicted negative | predicted positive |

## Predictive accuracy

- Accuracy is one of the measures we can compute based on the confusion matrix:


$$\text{Accuracy} = \frac{\text{TP} + \text{TN}}{\text{TP} + \text{TN} + \text{FP} + \text{FN}}$$

- In sklearn: use `confusion_matrix` and `accuracy_score` from `sklearn.metrics`.
- Accuracy is also the default evaluation measure for classification

```
confusion_matrix(y_test, y_pred):
[[49 4]
 [5 85]]
accuracy_score(y_test, y_pred): 0.9370629370629371
model.score(X_test, y_test): 0.9370629370629371
```

## The problem with accuracy: imbalanced datasets


- The type of error plays an even larger role if the dataset is imbalanced
  - One class is much more frequent than the other, e.g. credit fraud
  - Is a 99.99% accuracy good enough?
- Are these three models really equally good?



**Precision** is used when the goal is to limit FPs

- Clinical trials: you only want to test drugs that really work
- Search engines: you want to avoid bad search results


$$\text{Precision} = \frac{\text{TP}}{\text{TP} + \text{FP}}$$




**Recall** is used when the goal is to limit FNs

- Cancer diagnosis: you don't want to miss a serious disease
- Search engines: You don't want to omit important hits
- Also know as sensitivity, hit rate, true positive rate (TPR)

$$\text{Recall} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$



# Comparison



How many selected items are relevant?

$$\text{Precision} = \frac{\text{true positives}}{\text{true positives} + \text{false positives}}$$


How many relevant items are selected?

$$\text{Recall} = \frac{\text{true positives}}{\text{true positives} + \text{false negatives}}$$



**F1-score** or F1-measure trades off precision and recall:

$$F1 = 2 \cdot \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$



# Classification measure Zoo

|                     |                              | True condition                                                                                                                                      |                                                                                                                                              |                                                                                                                       |
|---------------------|------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------|
|                     |                              | Condition positive                                                                                                                                  | Condition negative                                                                                                                           |                                                                                                                       |
| Total population    |                              |                                                                                                                                                     |                                                                                                                                              | $\text{Prevalence} = \frac{\sum \text{Condition positive}}{\sum \text{Total population}}$                             |
| Predicted condition | Predicted condition positive | <b>True positive</b> ,<br>Power                                                                                                                     | <b>False positive</b> ,<br>Type I error                                                                                                      | $\text{Accuracy (ACC)} = \frac{\sum \text{True positive} + \sum \text{True negative}}{\sum \text{Total population}}$  |
|                     | Predicted condition negative | <b>False negative</b> ,<br>Type II error                                                                                                            | <b>True negative</b>                                                                                                                         | $\text{False discovery rate (FDR)} = \frac{\sum \text{False positive}}{\sum \text{Predicted condition positive}}$     |
|                     |                              | $\text{True positive rate (TPR), Recall, Sensitivity, probability of detection} = \frac{\sum \text{True positive}}{\sum \text{Condition positive}}$ | $\text{False positive rate (FPR), Fall-out, probability of false alarm} = \frac{\sum \text{False positive}}{\sum \text{Condition negative}}$ | $\text{Negative predictive value (NPV)} = \frac{\sum \text{True negative}}{\sum \text{Predicted condition negative}}$ |
|                     |                              | $\text{False negative rate (FNR), Miss rate} = \frac{\sum \text{False negative}}{\sum \text{Condition positive}}$                                   | $\text{Specificity (SPC), Selectivity, True negative rate (TNR)} = \frac{\sum \text{True negative}}{\sum \text{Condition negative}}$         | $\text{Diagnostic odds ratio (DOR)} = \frac{\text{LR+}}{\text{LR-}}$                                                  |
|                     |                              |                                                                                                                                                     | $\text{Positive likelihood ratio (LR+)} = \frac{\text{TPR}}{\text{FPR}}$                                                                     | $F_1 \text{ score} = \frac{1}{\frac{1}{\text{Recall}} + \frac{1}{\text{Precision}}}$                                  |
|                     |                              |                                                                                                                                                     | $\text{Negative likelihood ratio (LR-)} = \frac{\text{FNR}}{\text{TNR}}$                                                                     |                                                                                                                       |

[https://en.wikipedia.org/wiki/Precision and recall](https://en.wikipedia.org/wiki/Precision_and_recall)  
 (.[https://en.wikipedia.org/wiki/Precision and recall](https://en.wikipedia.org/wiki/Precision_and_recall))."


## Averaging scores per class

- Study the scores *by class* (in scikit-learn: `classification_report`)
  - One class viewed as positive, other(s) as negative
  - Support: number of samples in each class
  - Last line: weighted average over the classes (weighted by number of samples in each class)
- Averaging for scoring measure  $R$  across  $C$  classes (also for multiclass):
  - micro: count total number of TP, FP, TN, FN
  - macro

$$\frac{1}{C} \sum_{c \in C} R(y_c, \hat{y}_c)$$

- weighted ( $w_c$ : ratio of examples of class  $c$ )
$$\sum_{c \in C} w_c R(y_c, \hat{y}_c)$$

## Example

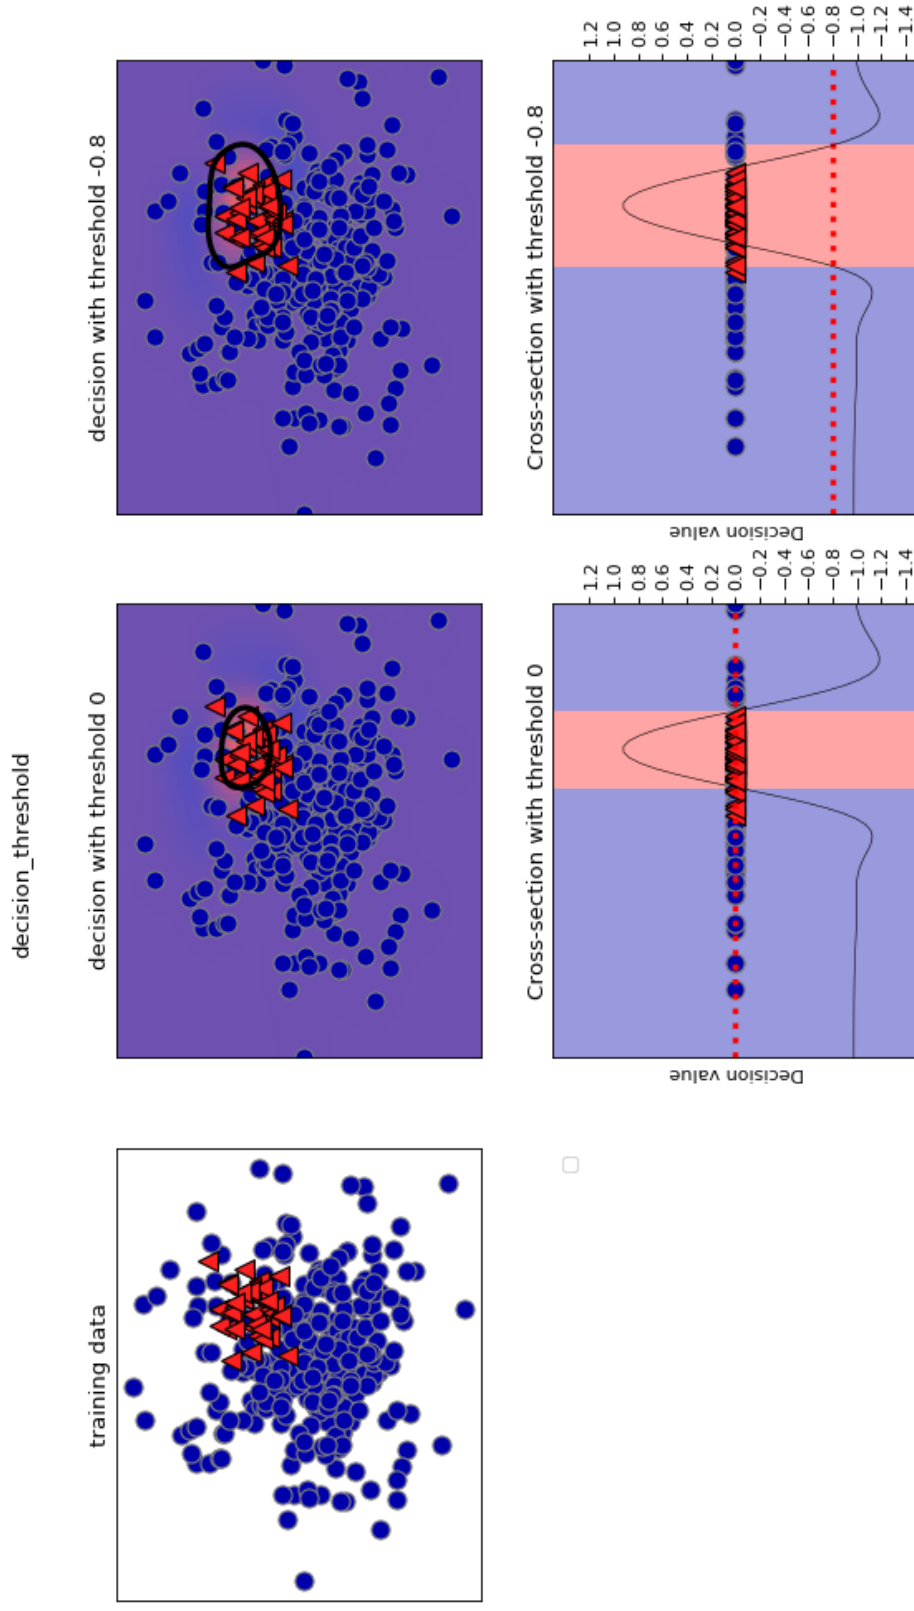


|          |              |           |        |          |         |
|----------|--------------|-----------|--------|----------|---------|
| Matrix 1 |              | precision | recall | f1-score | support |
|          | 0            | 0.90      | 1.00   | 0.95     | 90      |
|          | 1            | 0.00      | 0.00   | 0.00     | 10      |
|          | accuracy     |           |        | 0.90     | 100     |
|          | macro avg    | 0.45      | 0.50   | 0.47     | 100     |
|          | weighted avg | 0.81      | 0.90   | 0.85     | 100     |
| Matrix 2 |              | precision | recall | f1-score | support |
|          | 0            | 1.00      | 0.89   | 0.94     | 90      |
|          | 1            | 0.50      | 1.00   | 0.67     | 10      |
|          | accuracy     |           |        | 0.90     | 100     |
|          | macro avg    | 0.75      | 0.94   | 0.80     | 100     |
|          | weighted avg | 0.95      | 0.90   | 0.91     | 100     |
| Matrix 3 |              | precision | recall | f1-score | support |
|          | 0            | 0.94      | 0.94   | 0.94     | 90      |
|          | 1            | 0.50      | 0.50   | 0.50     | 10      |
|          | accuracy     |           |        | 0.90     | 100     |
|          | macro avg    | 0.72      | 0.72   | 0.72     | 100     |
|          | weighted avg | 0.90      | 0.90   | 0.90     | 100     |

# Taking uncertainty into account

- Remember that many classifiers actually return a probability per class
  - We can retrieve it with `decision_function` and `predict_proba`
- For binary classification, we threshold at 0 for `decision_function` and 0.5 for `predict_proba` by default
- However, depending on the evaluation measure, you may want to threshold differently to fit your goals
  - For instance, when a FP is much worse than a FN
  - This is called *threshold calibration*

- Imagine that we want to avoid misclassifying a positive (red) point
- Points within decision boundary (black line) are classified positive
- Lowering the decision threshold (bottom figure): fewer FN, more FP



- Studying the classification report, we see that lowering the threshold yields:
  - higher recall for class 1 (we risk more FPs in exchange for more TP)
  - lower precision for class 1
- We can often trade off precision for recall

| Threshold 0  |           |        |          |         |
|--------------|-----------|--------|----------|---------|
|              | precision | recall | f1-score | support |
| 0            | 0.91      | 0.96   | 0.93     | 96      |
| 1            | 0.67      | 0.47   | 0.55     | 17      |
| accuracy     |           |        | 0.88     | 113     |
| macro avg    | 0.79      | 0.71   | 0.74     | 113     |
| weighted avg | 0.87      | 0.88   | 0.88     | 113     |

| Threshold -0.8 |           |        |          |         |
|----------------|-----------|--------|----------|---------|
|                | precision | recall | f1-score | support |
| 0              | 0.98      | 0.92   | 0.95     | 96      |
| 1              | 0.65      | 0.88   | 0.75     | 17      |
| accuracy       |           |        | 0.91     | 113     |
| macro avg      | 0.81      | 0.90   | 0.85     | 113     |
| weighted avg   | 0.93      | 0.91   | 0.92     | 113     |

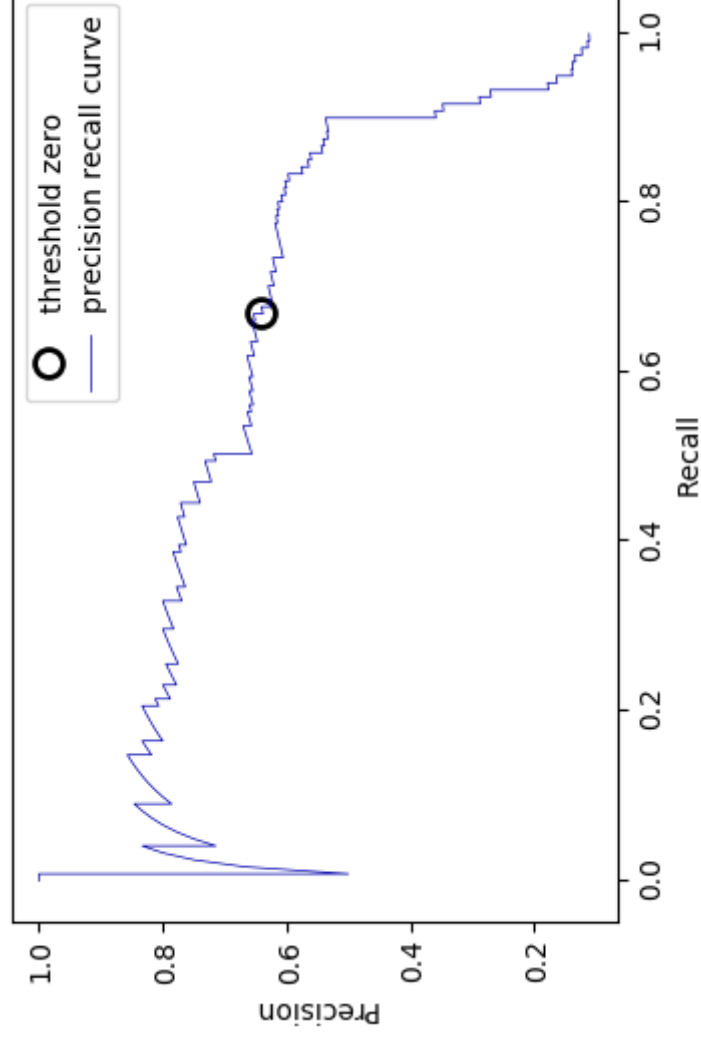


## Precision-Recall curves

- The best threshold depends on your application, should be driven by real-world goals.
- You can have arbitrary high recall, but you often want reasonable precision, too.
- It is not clear beforehand where the optimale trade-off (or *operating point*) will be, so it is useful to look at all possible thresholds
- Plotting precision against recall for all thresholds yields a **precision-recall curve**


- In scikit-learn, this is included in the `sklearn.metrics` module
- Returns all precision and recall values for all thresholds
  - Vary threshold from lowest to highest decision function score in the predictions
  - Or from highest to lowest class probability

- The default tradeoff (chosen by the predict method) is shown as *threshold zero*.
  - Higher threshold, more precision (move left)
  - Lower threshold, more recall (move right)
- The closer the curve stays to the upper-right corner, the better
  - High precision and high recall
- Here, it is possible to still get a precision of 0.5 with high recall



## Model selection

- Different classifiers work best in different parts of the curve (at different operating points)
- RandomForest (in red) performs better at the extremes, SVM better in center
- The area under the precision-recall curve (AUPRC) is often used as a general evaluation measure



Note that the F1-measure completely misses these subtleties

```
f1_score of random forest: 0.610
f1_score of svc: 0.656
```

- The area under the precision-recall curve is returned by the `average_precision_score` measure
  - It's actually a close approximation of the actual area
- This is a good automatic measure, but also hides the subtleties

Average precision of random forest: 0.660

Average precision of svc: 0.666


# Receiver Operating Characteristics (ROC) and AUC

- There is another trade-off between recall (true positive rate, TPR) and the false positive rate (FPR).
- The 2D space created by TPR and FPR is called the Receiver Operating Characteristics (ROC) space
- A model will be at one point in this ROC space

$$\text{TPR} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$

$$\text{FPR} = \frac{\text{FP}}{\text{FP} + \text{TN}}$$


- Varying the decision threshold yields the ROC curve
- It can be computed with the `roc_curve` function
  - Lower threshold, more recall/TPR, move right
  - High threshold, fewer FPs, move left
- Ideal is close to the top left: high recall, low FPR
- Inspect the curve to find the preferred calibration
  - Here, we can get much higher recall with slightly worse FPR






## Visualization

- The blue probability density shows the probability  $p(x)$  that the model predicts blue if a data point has a certain predicted probability  $x$  to be blue. Same for red.
- In a random classifier the probability densities completely overlap.
- All points with a predicted probability higher than the threshold are predicted positive, others negative
- As we increase the threshold, we'll get fewer FPs, more FNs. We move from right to left along the ROC curve.




|    |    |
|----|----|
| TP | FP |
| FN | TN |
| 1  | 1  |




## ROC Isometrics

- Different *costs* can be involved for FP and FN
- This yields different *isometrics* (lines of equal cost) in ROC space
- The optimal threshold is the point on the ROC curve where the cost is minimal
  - If a FP and FN are weighed equally, cost lines follow the diagonal (blue line)
  - If a FP is 10 times worse than a FN: pink line
  - If a FN is 10 times worse than a FP: red line



## Model selection

- Again, we can compare multiple models by looking at the ROC curves
- We can calibrate the threshold depending on whether we need high recall or low FPR
- We can select between algorithms (or hyperparameters) depending on the involved costs.



## Area under the ROC curve

- A good summary measure is the area under the ROC curve (AUROC or AUC)
- Compute using the `roc_auc_score`
  - Don't use `auc` (uses less accurate trapezoidal rule)

```
rf_auc = roc_auc_score(y_test, rf.predict_proba(X_test)[: , 1])
svc_auc = roc_auc_score(y_test, svc.decision_function(X_test))
```

AUC for Random Forest: 0.937


AUC for SVC: 0.916

## Imbalanced classes

- AUC is popular because it is insensitive to class imbalance
  - Random guessing always yields  $TPR=FPR$
  - All points are on the diagonal line, hence an AUC of 0.5
  - Hint: use the visualization of TPR, FPR to see this

- Example: unbalanced digits
  - 3 models, ACC is the same, AUC not
  - If we optimize for ACC, our model could be just random guessing

|               |                 |              |
|---------------|-----------------|--------------|
| gamma = 1.000 | accuracy = 0.90 | AUC = 0.5000 |
| gamma = 0.100 | accuracy = 0.90 | AUC = 0.9582 |
| gamma = 0.010 | accuracy = 0.90 | AUC = 0.9995 |



## **Take home message**

- AUC is highly recommended, especially on imbalanced data
- Remember to calibrate the threshold to your needs



# Multi-class classification

- Multiclass metrics are derived from binary metrics, averaged over all classes
- Let's consider the full (10-class) handwritten digit recognition data


Confusion matrix

Accuracy: 0.953  
Confusion matrix:

|      |    |    |    |    |    |    |    |    |     |
|------|----|----|----|----|----|----|----|----|-----|
| [ 37 | 0  | 0  | 0  | 0  | 0  | 0  | 0  | 0  | 0]  |
| [ 0  | 39 | 0  | 0  | 0  | 0  | 2  | 0  | 2  | 0]  |
| [ 0  | 0  | 41 | 3  | 0  | 0  | 0  | 0  | 0  | 0]  |
| [ 0  | 0  | 1  | 43 | 0  | 0  | 0  | 0  | 0  | 1]  |
| [ 0  | 0  | 0  | 0  | 38 | 0  | 0  | 0  | 0  | 0]  |
| [ 0  | 1  | 0  | 0  | 0  | 47 | 0  | 0  | 0  | 0]  |
| [ 0  | 0  | 0  | 0  | 0  | 0  | 52 | 0  | 0  | 0]  |
| [ 0  | 1  | 0  | 1  | 1  | 0  | 0  | 45 | 0  | 0]  |
| [ 0  | 3  | 1  | 0  | 0  | 0  | 0  | 0  | 43 | 1]  |
| [ 0  | 0  | 0  | 1  | 0  | 1  | 0  | 0  | 1  | 44] |

## Visualized as a heatmap

- Which digits are easy to predict? Which ones are confused?



Precision, recall, F1-score now yield 10 per-class scores

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 1.00      | 1.00   | 1.00     | 37      |
| 1            | 0.89      | 0.91   | 0.90     | 43      |
| 2            | 0.95      | 0.93   | 0.94     | 44      |
| 3            | 0.90      | 0.96   | 0.92     | 45      |
| 4            | 0.97      | 1.00   | 0.99     | 38      |
| 5            | 0.98      | 0.98   | 0.98     | 48      |
| 6            | 0.96      | 1.00   | 0.98     | 52      |
| 7            | 1.00      | 0.94   | 0.97     | 48      |
| 8            | 0.93      | 0.90   | 0.91     | 48      |
| 9            | 0.96      | 0.94   | 0.95     | 47      |
| accuracy     |           |        | 0.95     | 450     |
| macro avg    | 0.95      | 0.95   | 0.95     | 450     |
| weighted avg | 0.95      | 0.95   | 0.95     | 450     |

## Different ways to compute average

- macro-averaging: computes unweighted per-class scores:  $\frac{\sum_{i=0}^n score_i}{n}$ 
  - Use when you care about each class equally much
- weighted averaging: scores are weighted by the relative size of the classes (support):  $\frac{\sum_{i=0}^n score_i weight_i}{n}$ 
  - Use when data is imbalanced
- micro-averaging: computes total number of FP, FN, TP over all classes, then computes scores using these counts:  $recall = \frac{\sum_{i=0}^n TP_i}{\sum_{i=0}^n TP_i + \sum_{i=0}^n FN_i}$ 
  - Use when you care about each sample equally much

Micro average f1 score: 0.953

Weighted average f1 score: 0.953



Macro average f1 score: 0.954

# Regression metrics

Most commonly used are


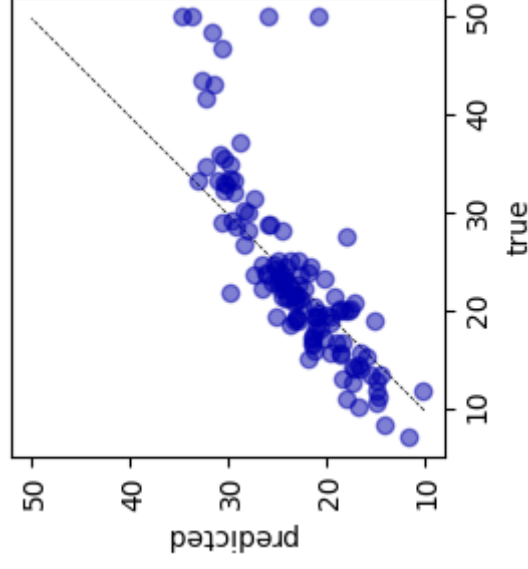
- (root) mean squared error:  $\frac{\sum_i (y_{pred_i} - y_{actual_i})^2}{n}$
- mean absolute error:  $\frac{\sum_i |y_{pred_i} - y_{actual_i}|}{n}$ 
  - Less sensitive to outliers and large errors
- R squared (r2):  $1 - \frac{\sum_i (y_{pred_i} - y_{actual_i})^2}{\sum_i (y_{mean} - y_{actual_i})^2}$ 
  - Ratio of variation explained by the model / total variation
  - Between 0 and 1, but *negative* if the model is worse than just predicting the mean
  - Easier to interpret (higher is better).

- R squared: 1 - ratio of  $\sum_i (y_{pred_i} - y_{actual_i})^2$  (blue) and  $\sum_i (y_{mean} - y_{actual_i})^2$  (red)



## Visualizing errors

- Prediction plot (left): predicted vs actual target values
- Residual plot (right): residuals vs actual target values
  - Over- and underpredictions can be given different costs



# Using evaluation metrics in model selection

- You typically want to use AUC or other relevant measures in `cross_val_score` and `GridSearchCV` instead of the default accuracy.
- scikit-learn makes this easy through the `scoring` argument
  - But, you need to need to look the mapping between the scorer and the metric ([http://scikit-learn.org/stable/modules/model\\_evaluation.html#model-evaluation](http://scikit-learn.org/stable/modules/model_evaluation.html#model-evaluation)).



| Scoring                     | Function                        | Comment                        |
|-----------------------------|---------------------------------|--------------------------------|
| <b>Classification</b>       |                                 |                                |
| 'accuracy'                  | metrics.accuracy_score          |                                |
| 'average_precision'         | metrics.average_precision_score |                                |
| 'f1'                        | metrics.f1_score                | for binary targets             |
| 'f1_micro'                  | metrics.f1_score                | micro-averaged                 |
| 'f1_macro'                  | metrics.f1_score                | macro-averaged                 |
| 'f1_weighted'               | metrics.f1_score                | weighted average               |
| 'f1_samples'                | metrics.f1_score                | by multilabel sample           |
| 'neg_log_loss'              | metrics.log_loss                | requires predict_proba support |
| 'precision' etc.            | metrics.precision_score         | suffixes apply as with 'f1'    |
| 'recall' etc.               | metrics.recall_score            | suffixes apply as with 'f1'    |
| 'roc_auc'                   | metrics.roc_auc_score           |                                |
| <b>Clustering</b>           |                                 |                                |
| 'adjusted_rand_score'       | metrics.adjusted_rand_score     |                                |
| <b>Regression</b>           |                                 |                                |
| 'neg_mean_absolute_error'   | metrics.mean_absolute_error     |                                |
| 'neg_mean_squared_error'    | metrics.mean_squared_error      |                                |
| 'neg_median_absolute_error' | metrics.median_absolute_error   |                                |
| 'r2'                        | metrics.r2_score                |                                |

Or simply look up like this:

Available scorers:

```
['accuracy', 'adjusted_mutual_info_score', 'adjusted_rand_score', 'average_
_precision', 'balanced_accuracy', 'brier_score_loss', 'completeness_scor
e', 'explained_variance', 'f1', 'f1_macro', 'f1_micro', 'f1_samples', 'f1_
weighted', 'fowlkes_mallows_score', 'homogeneity_score', 'jaccard', 'jacc
ard_macro', 'jaccard_micro', 'jaccard_samples', 'jaccard_weighted', 'max_er
ror', 'mutual_info_score', 'neg_log_loss', 'neg_mean_absolute_error', 'neg
_mean_squared_error', 'neg_mean_squared_log_error', 'neg_median_absolute_e
rror', 'normalized_mutual_info_score', 'precision', 'precision_macro', 'pr
ecision_micro', 'precision_samples', 'precision_weighted', 'r2', 'recall',
'recall_macro', 'recall_micro', 'recall_samples', 'recall_weighted', 'roc_
auc', 'v_measure_score']
```

## Cross-validation with AUC

```
roc_auc = cross_val_score(SVC(), digits.data, digits.target == 9,
 scoring="roc_auc")
print("AUC scoring: {}".format(roc_auc))
```

```
Default scoring: [0.9 0.9 0.9]
Explicit accuracy scoring: [0.9 0.9 0.9]
AUC scoring: [0.994 0.99 0.996]
```

## Grid Search with accuracy and AUC

- With accuracy, gamma=0.0001 is selected
- With AUC, gamma=0.01 is selected
  - Actually has better accuracy on the test set

```
Grid-Search with accuracy
Best parameters: {'gamma': 0.0001}
Best cross-validation score (accuracy): 0.970
Test set AUC: 0.992
Test set accuracy: 0.973
```

```
Grid-Search with AUC
Best parameters: {'gamma': 0.01}
Best cross-validation score (AUC): 0.997
Test set AUC: 1.000
Test set accuracy: 1.000
```

# Final thoughts

- There exist techniques to correct label imbalance
  - Undersample the majority class, or oversample the minority class
  - SMOTE (Synthetic Minority Oversampling TEchnique) adds artificial *training* points by interpolating existing minority class points
    - Think twice before creating 'artificial' training data
- Cost-sensitive classification (not in sklearn)
  - *Cost matrix*: a confusion matrix with a costs associated to every possible type of error
  - Some algorithms allow optimizing on these costs instead of their usual loss function
  - Meta-cost: builds ensemble of models by relabeling training sets to match a given cost matrix
    - Black-box: can make any algorithm cost sensitive (but slower and less accurate)

- There are many more metrics to choose from
  - Cohen's Kappa: accuracy, taking into account the possibility of predicting the right class by chance
    - 1: perfect prediction, 0: random prediction, negative: worse than random
    - With  $p_0$  = accuracy, and  $p_e$  = accuracy of random classifier:
 
$$\kappa = \frac{p_o - p_e}{1 - p_e}$$
  - Balanced accuracy: accuracy where each sample is weighted according to the inverse prevalence of its true class
    - Identical to macro-averaged recall
  - Matthews correlation coefficient: another measure that can be used on imbalanced data
    - 1: perfect prediction, 0: random prediction, -1: inverse prediction

$$MCC = \frac{tp \times tn - fp \times fn}{\sqrt{(tp + fp)(tp + fn)(tn + fp)(tn + fn)}}$$

# Bias-Variance decomposition

- When we repeat evaluation procedures multiple times, we can distinguish two sources of errors:
  - Bias: systematic error (independent of the training sample). The classifier always gets certain points wrong
  - Variance: error due to variability of the model with respect to the training sample. The classifier predicts some points accurately on some training sets, but inaccurately on others.
- There is also an intrinsic (noise) error, but there's nothing we can do against that.
- Bias is associated with underfitting, and variance with overfitting
- Bias-variance trade-off: you can often exchange variance for bias through regularization (and vice versa)
  - The challenge is to find the right trade-off (minimizing total error)
- Useful to understand how to tune or adapt learning algorithm

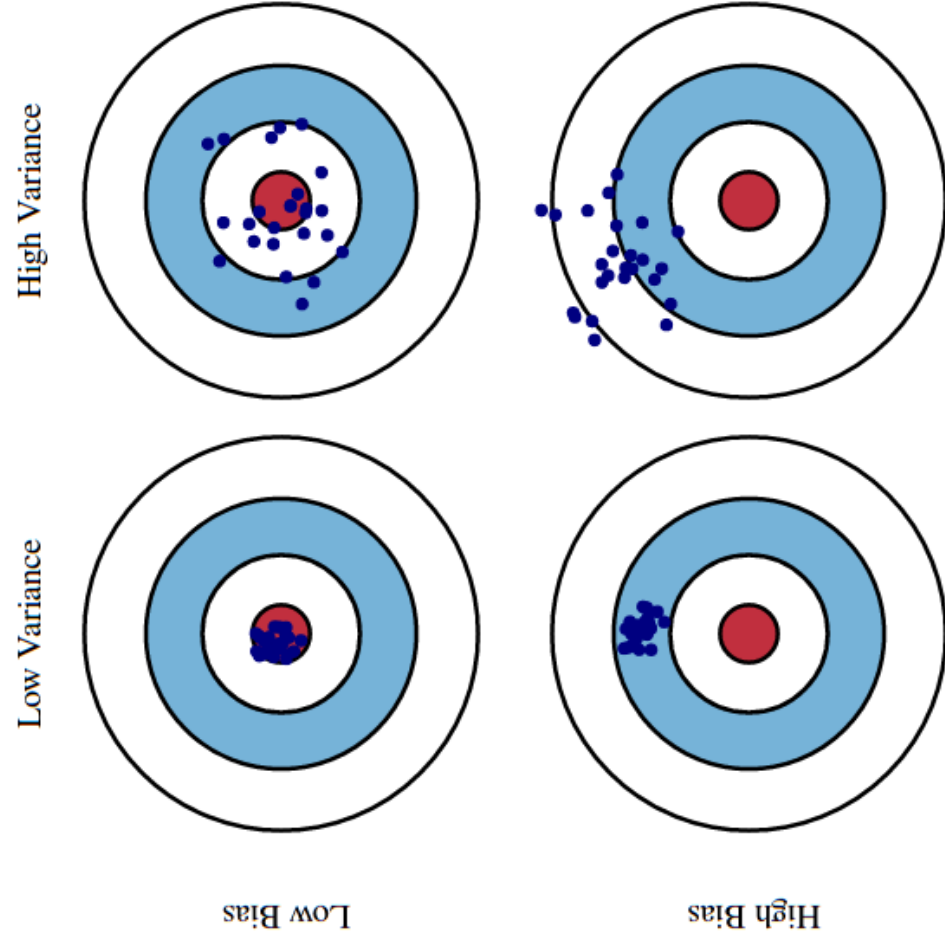


Fig. 1 Graphical illustration of bias and variance.



- Sadly, this is not yet supported by scikit-learn
- How to measure bias and variance (for regression):
  - Take 100 or more bootstraps (or shuffle-splits)
  - For each data point  $x$ :
    - $bias(x)^2 = (x_{true} - mean(x_{predicted}))^2$
    - $variance(x) = var(x_{predicted})$
  - Total bias:  $\sum_x bias(x)^2 * w_x$ , with  $w_x$  the ratio of  $x$  occurring in the test set
  - Total variance:  $\sum_x variance(x) * w_x$

- General procedure for (binary) classification:
  - Take 100 or more bootstraps (or shuffle-splits)
  - Bias for any point  $x$  = misclassification ratio
    - If misclassified 50% of the time:  $bias(x) = 0.5$
  - Variance for any point  $x$  is  $(1 - (P(class_1)^2 + P(class_2)^2))/2$ 
    - $P(class_i)$  is ratio of class  $i$  predictions
    - When each class predicted half of the time:  
 $variance(x) = (1 - (0.5^2 + 0.5^2))/2 = 0.25$
  - Total bias:  $\sum_x bias(x)^2 * w_x$ , with  $w_x$  the ratio of  $x$  occurring in the test data
  - Total variance:  $\sum_x variance(x) * w_x$

```


for i, (train_index, test_index) in enumerate(shuffle_split.split(X)):
 clf.fit(X[train_index], y[train_index])
 y_pred = clf.predict(X[test_index])
 # Store predictions
 for i, index in enumerate(test_index):
 y_all_pred[index].append(y_pred[i])

Compute bias, variance, error
bias_sq = sum([(1 - x.count(y[i])/len(x))**2 * len(x)/n_repeat
 for i, x in enumerate(y_all_pred)])
var = sum([(1 - ((x.count(0)/len(x))**2 + (x.count(1)/len(x))**2))/2) *
 len(x)/n_repeat
 for i, x in enumerate(y_all_pred)])
error = sum([(1 - x.count(y[i])/len(x)) * len(x)/n_repeat
 for i, x in enumerate(y_all_pred)])

```


Bias squared: 14.50, Variance: 0.84, Total error: 15.34

# Bias-variance and overfitting



- High bias means that you are likely underfitting
  - Do less regularization
  - Use a more flexible/complex model (another algorithm)
  - Use a bias-reduction technique (e.g. boosting, see later)
- High variance means that you are likely overfitting
  - Use more regularization
  - Get more data
  - Use a simpler model (another algorithm)
  - Use a variance-reduction techniques (e.g. bagging, see later)

## Bias-Variance Flowchart (Andrew Ng, Coursera)



# Hyperparameter tuning

Now that we know how to evaluate models, we can improve them by tuning their hyperparameters

We can basically use any optimization technique to optimize hyperparameters:

- **Grid search**
- **Random search**

More advanced techniques:

- Local search
- Racing algorithms
- Model-based optimization (see later)
- Multi-armed bandits
- Genetic algorithms



# Grid Search

- For each hyperparameter, create a list of interesting/possible values
  - E.g. For kNN: k in [1,3,5,7,9,11,33,55,77,99]
- Evaluate all possible combination of hyperparameter values
  - E.g. using cross-validation
- Select the hyperparameter values yielding the best results
- A naive approach would be to just loop over all combinations

```
for gamma in [0.001, 0.01, 0.1, 1, 10, 100]:
 for C in [0.001, 0.01, 0.1, 1, 10, 100]:
 # for each combination, train and evaluate an SVC
 svm = SVC(gamma=gamma, C=C);
 svm.fit(X_train, y_train);
 score = svm.score(X_test, y_test)
 # if we got a better score, store the score and parameters
 if score > best_score:
 best_score = score
 best_parameters = {'C': C, 'gamma': gamma}
```


Size of training set: 112    size of test set: 38

Best score: 0.97

Best parameters: {'C': 100, 'gamma': 0.001}

## Overfitting the parameters and the validation set

- Simply taking the best performing model yields optimistic results
- We've already used the test data to evaluate each hyperparameter setting!
- Hence, we don't have an independent test set to evaluate these hyperparameter settings
  - Information 'leaks' from test set into the final model
- Solution: Set aside part of the training data to evaluate the hyperparameter settings
  - Select best hyperparameters on validation set
  - Rebuild the model on the training+validation set
  - Evaluate optimal model on the test set



```
split data into train+validation set and test set
X_trainval, X_test, y_trainval, y_test = train_test_split(
 iris.data, iris.target, random_state=0)
split train+validation set into training and validation set
X_train, X_valid, y_train, y_valid = train_test_split(
 X_trainval, y_trainval, random_state=1)
```

Size of training set: 84    size of validation set: 28    size of test set:  
38

Best score on validation set: 0.96  
Best parameters: {'C': 10, 'gamma': 0.001}  
Test set score with best parameters: 0.92

## **Grid-search with cross-validation**

- Same problem as before: the way that we split the data into training, validation, and test set may have a large influence on estimated performance
- We need to use cross-validation again, instead of a single split
- Expensive. Often, 3 or 5-fold CV is enough

```

for gamma in [0.001, 0.01, 0.1, 1, 10, 100]:
 for C in [0.001, 0.01, 0.1, 1, 10, 100]:
 # train an SVC
 svm = SVC(gamma=gamma, C=C)
 # perform cross-validation
 scores = cross_val_score(svm, X_trainval, y_trainval, cv=5)
 # compute mean cross-validation accuracy
 score = np.mean(scores)
 # if we got a better score, store the score and parameters
 if score > best_score:
 best_score = score
 best_parameters = {'C': C, 'gamma': gamma}

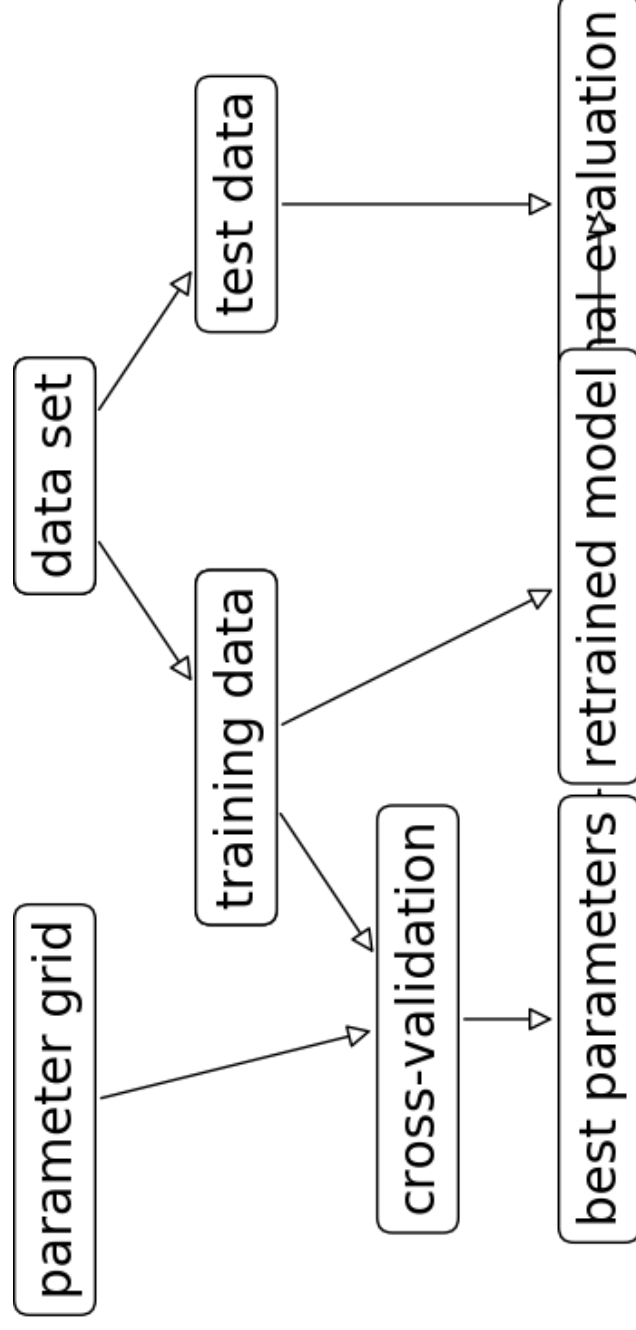
```

```

Out[51]: SVC(C=100, cache_size=200, class_weight=None, coef0=0.0,
 decision_function_shape='ovr', degree=3, gamma=0.01, kernel='rbf',
 max_iter=-1, probability=False, random_state=None, shrinking=True,
 tol=0.001, verbose=False)

```

## Overall process



## Grid search in scikit-learn

- Create a parameter grid as a dictionary
  - Keys are parameter names
  - Values are lists of hyperparameter values

```
param_grid = {'C': [0.001, 0.01, 0.1, 1, 10, 100],
 'gamma': [0.001, 0.01, 0.1, 1, 10, 100]}
print("Parameter grid:\n{}".format(param_grid))
```

```
Parameter grid:
{'C': [0.001, 0.01, 0.1, 1, 10, 100], 'gamma': [0.001, 0.01, 0.1, 1, 10, 1
00]}
```



- GridSearchCV: like a classifier that uses CV to automatically optimize its hyperparameters internally
  - Input: (untrained) model, parameter grid, CV procedure
  - Output: optimized model on given training data
  - Should only have access to training data

```
grid_search = GridSearchCV(SVC(), param_grid, cv=5)
grid_search.fit(X_train, y_train)
```

Out[56]:

```
GridSearchCV(cv=5, error_score='raise-deprecating',
 estimator=SVC(C=1.0, cache_size=200, class_weight=None, coef
0=0.0,
 decision_function_shape='ovr', degree=3,
 gamma='auto_deprecated', kernel='rbf', max_iter
r=-1,
 probability=False, random_state=None, shrinkin
g=True,
 tol=0.001, verbose=False),
 iid='warn', n_jobs=None,
 param_grid={'C': [0.001, 0.01, 0.1, 1, 10, 100],
 'gamma': [0.001, 0.01, 0.1, 1, 10, 100]},
 pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
 scoring=None, verbose=0)
```

The optimized test score and hyperparameters can easily be retrieved:

```
grid_search.score(X_test, y_test)
```

```
Test set score: 0.97
```

```
grid_search.best_params_
grid_search.best_score_
```

```
Best parameters: {'C': 100, 'gamma': 0.01}
Best cross-validation score: 0.97
```

```
grid_search.best_estimator_
```

```
Best estimator:
SVC(C=100, cache_size=200, class_weight=None, coef0=0.0,
 decision_function_shape='ovr', degree=3, gamma=0.01, kernel='rbf',
 max_iter=-1, probability=False, random_state=None, shrinking=True,
 tol=0.001, verbose=False)
```

## Visualizing hyperparameter impact


We can retrieve and visualize the cross-validation result to better understand the impact of hyperparameters

```
results = pd.DataFrame(grid_search.cv_results_)
```

|   | mean_fit_time | std_fit_time | mean_score_time | std_score_time | ... | split4_test_score | mean_test_score | std_test_score | rank_test_score |
|---|---------------|--------------|-----------------|----------------|-----|-------------------|-----------------|----------------|-----------------|
| 0 | 1.06e-03      | 5.41e-04     | 4.11e-04        | 1.21e-04       | ... | 0.38              | 0.37            | 0.01           | 22              |
| 1 | 5.87e-04      | 5.28e-05     | 2.51e-04        | 7.88e-06       | ... | 0.38              | 0.37            | 0.01           | 22              |
| 2 | 6.44e-04      | 8.58e-05     | 3.07e-04        | 1.02e-04       | ... | 0.38              | 0.37            | 0.01           | 22              |
| 3 | 8.14e-04      | 1.50e-04     | 3.85e-04        | 6.89e-05       | ... | 0.38              | 0.37            | 0.01           | 22              |
| 4 | 9.86e-04      | 2.02e-04     | 4.03e-04        | 1.13e-04       | ... | 0.38              | 0.37            | 0.01           | 22              |

5 rows × 15 columns

Visualize as a heatmap



When hyperparameters depend on other parameters, we can use lists of dictionaries to define the hyperparameter space

```
param_grid = [{'kernel': ['rbf'],
 'C': [0.001, 0.01, 0.1, 1, 10, 100],
 'gamma': [0.001, 0.01, 0.1, 1, 10, 100]},
 {'kernel': ['linear'],
 'C': [0.001, 0.01, 0.1, 1, 10, 100]}]
```

List of grids:

```
[{'kernel': ['rbf'], 'C': [0.001, 0.01, 0.1, 1, 10, 100], 'gamma': [0.001,
0.01, 0.1, 1, 10, 100]}, {'kernel': ['linear'], 'C': [0.001, 0.01, 1,
10, 100]}]
```

## Nested cross-validation

- Note that we are still using a single split to create the outer test set
- We can also use cross-validation here
- Nested cross-validation:
  - Outer loop: split data in training and test sets
  - Inner loop: run grid search, splitting the training data into train and validation sets
- Result is a just a list of scores
  - There will be multiple optimized models and hyperparameter settings (not returned)
- To apply on future data, we need to train GridSearchCV on all data again

```
scores = cross_val_score(GridSearchCV(SVC(), param_grid, cv=5),
 iris.data, iris.target, cv=5)
```

```
Cross-validation scores: [0.967 1. 0.9 0.967 1.]
Mean cross-validation score: 0.966666666666666668
```

## Parallelizing cross-validation and grid-search

- On a practical note, it is easy to parallelize CV and grid search
- `cross_val_score` and `GridSearchCV` have a `n_jobs` parameter defining the number of cores it can use.
  - set it to `n_jobs=-1` to use all available cores.



# Random Search

- Grid Search has a few downsides:
  - Optimizing many hyperparameters creates a combinatorial explosion
  - You have to predefine a grid, hence you may jump over optimal values
- Random Search:
  - Picks `n_iter` random parameter values
  - Scales better, you control the number of iterations
  - Often works better in practice, too
    - not all hyperparameters interact strongly
    - you don't need to explore all combinations

- Executing random search in scikit-learn:
  - RandomizedSearchCV works like GridSearchCV
  - Has `n_iter` parameter for the number of iterations
  - Search grid can use distributions instead of fixed lists

```
param_grid = {'C': expon(scale=100),
 'gamma': expon(scale=.1)}
random_search = RandomizedSearchCV(SVC(), param_distributions=param_grid
,
 n_iter=20)
random_search.fit(X_train, y_train)
```

Out[64]:

```
RandomizedSearchCV(cv='warn', error_score='raise-deprecating',
 estimator=SVC(C=1.0, cache_size=200, class_weight=None
e,
 r',
 coef0=0.0, decision_function_shape='ov
 degree=3, gamma='auto_deprecated',
 kernel='rbf', max_iter=-1, probability=F
 else,
 random_state=None, shrinking=True, tol=
 0.001,
 verbose=False),
 iid='warn', n_iter=20, n_jobs=None,
 param_distributions={'C': <scipy.stats._distn_infrastr
ucture.rv_frozen object at 0x11d861668>,
 'gamma': <scipy.stats._distn_infr
astructure.rv_frozen object at 0x11d98a860>},
 pre_dispatch='2*n_jobs', random_state=None, refit=True
e,
 return_train_score=False, scoring=None, verbose=0)
```

# Summary

- k-fold Cross-validation
  - Choose k depending on how much data you have
    - Larger k is slower, but allows more training data
    - 10-fold, 5-fold, 5x2-fold most popular
  - Always use stratification for (imbalanced) classification
  - Train-test split and Shuffle-split: useful for large datasets
  - Use grouping when you want to generalize over groups
- Model selection
  - Don't aggregate over test scores: those have seen the test data
  - Use validation sets to choose algorithms/hyperparameters first
- Optimization
  - Grid Search: exhaustive but simple
  - Random Search: scales better
  - We'll see more advanced techniques later

# Summary

- Real-world data is often imbalanced
- False positives may be much worse than false negatives (or vice-versa)
- Binary classification
  - Select metrics that can distinguish different types of errors (precision, recall, f1-score, AUC,...)
  - Calibrate decision thresholds to the task at hand
  - Precision-Recall and ROC curves: choose the best threshold or take area under the curve
- Multiclass classification
  - Macro/Micro/weighted average of per-class scores (one-vs-all)
- Regression
  - (Root) mean squared/absolute error from 0..Inf
  - R2 easier to interpret
- All measures can be used in cross-validation or grid/random search
- Cost-sensitive classification: optimize for any cost matrix or cost function