


# Trees and Ensembles

# Trees



## Building Decision Trees


- Split the data in two (or more) parts
- Search over all possible splits and choose the one that is most *informative*
  - Many heuristics
  - E.g. *information gain*: how much does the entropy of the class labels decrease after the split (purer 'leafs')
- Repeat recursive partitioning



Making predictions:



- Classification: find leaf for new data point, predict majority class (or class distribution)
- Regression: idem, but predict the *mean* of all values


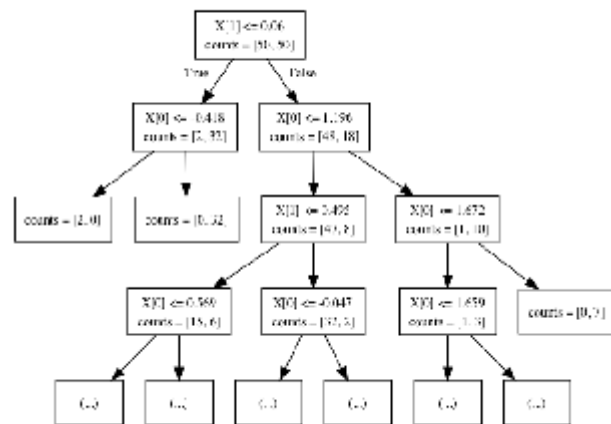
## Decision Tree classification

Where would you make the first split?









# Heuristics

For classification ( $X_i \rightarrow class_k$ ): *Impurity measures*:

- Misclassification Error (leads to larger trees):

$$1 - \operatorname{argmax}_k \hat{p}_k$$

- Gini-Index (probabilistic predictions):

$$\sum_{k \neq k'} \hat{p}_k \hat{p}_{k'} = \sum_{k=1}^K \hat{p}_k (1 - \hat{p}_k)$$

with  $\hat{p}_k$  = the relative frequency of class  $k$  in the leaf node



- Entropy (of the class attribute) measures *unpredictability* of the data:
  - How likely will random example have class  $k$ ?

$$E(X) = - \sum_{k=1}^K \hat{p}_k \log_2 \hat{p}_k$$


- Information Gain (a.k.a. Kullback–Leibler divergence) for choosing attribute  $X_i$  to split the data:

$$G(X, X_i) = E(X) - \sum_{v=1}^V \frac{|X_{i=v}|}{|X_i|} E(X_{i=v})$$

with  $\hat{p}_k$  = the relative frequency of class  $k$  in the leaf node,  $X$  = the training set, containing  $i$  features (variables)  $X_i$ ,  $v$  a specific value for  $X_i$ ,  $X_{i=v}$  is the set of examples having value  $v$  for feature  $X_i$ :  $\{x \in X | X_i = v\}$

## Heuristics visualized (binary class)

- Note that  $\text{gini} \neq \text{entropy}/2$



## Example

Ex.	1	2	3	4	5	6
a1	T	T	T	F	F	F
a2	T	T	F	F	T	T
class	+	+	-	+	-	-

$E(X)$  ?

$G(X, X_{a2})$  ?

$G(X, X_{a1})$  ?

$E(X) = -(\frac{1}{2} * \log_2(\frac{1}{2}) + \frac{1}{2} * \log_2(\frac{1}{2})) = 1$  (classes have equal probabilities)

$G(X, X_{a2}) = 0$  (after split, classes still have equal probabilities, entropy stays 1)

Ex.	1	2	3	4	5	6
a1	T	T	T	F	F	F
a2	T	T	F	F	T	T
class	+	+	-	+	-	-

$$E(X) = - \sum_{k=1}^K \hat{p}_k \log \hat{p}_k \quad , \quad G(X, X_i) = E(X) - \sum_{v=1}^V \frac{|X_{i=v}|}{|X_i|} E(X_{i=v})$$

$$E(X_{a1=T}) = -\frac{2}{3} \log_2\left(\frac{2}{3}\right) - \frac{1}{3} \log_2\left(\frac{1}{3}\right) = 0.9183 \quad (= E(X_{a1=F}))$$

$$G(X, X_{a1}) = 1 - \frac{1}{2} 0.9183 - \frac{1}{2} 0.9183 = 0.0817$$










hence we split on a1

## Heuristics in scikit-learn

The splitting criterion can be set with the `criterion` option in `DecisionTreeClassifier`

- `gini` (default): gini impurity index
- `entropy`: information gain

Best value depends on dataset, as well as other hyperparameters



# Handling many-valued features

What happens when a feature has (almost) as many values as examples?

- Information Gain will select it

One approach: use Gain Ratio instead (not available scikit-learn):

$$\text{GainRatio}(X, X_i) = \frac{\text{Gain}(X, X_i)}{\text{SplitInfo}(X, X_i)}$$

$$\text{SplitInfo}(X, X_i) = - \sum_{v=1}^V \frac{|X_{i=v}|}{|X|} \log_2 \frac{|X_{i=v}|}{|X|}$$

where  $X_{i=v}$  is the subset of examples for which feature  $X_i$  has value  $v$ .

SplitInfo will be big if  $X_i$  fragments the data into many small subsets, resulting in a smaller Gain Ratio.

# Overfitting: Controlling complexity of Decision Trees

Decision trees can very easily overfit the data. Regularization strategies:

- Pre-pruning: stop creation of new leafs at some point
  - Limiting the depth of the tree, or the number of leafs
  - Requiring a minimal leaf size (number of instances)
- Post-pruning: build full tree, then prune (join) leafs
  - Reduced error pruning: evaluate against held-out data
  - Many other strategies exist.
  - scikit-learn supports none of them (yet)



Effect of pre-pruning: default tree overfits, setting max\_depth=5 is better

```
tree = DecisionTreeClassifier(random_state=0)
tree.fit(X_train, y_train)
```

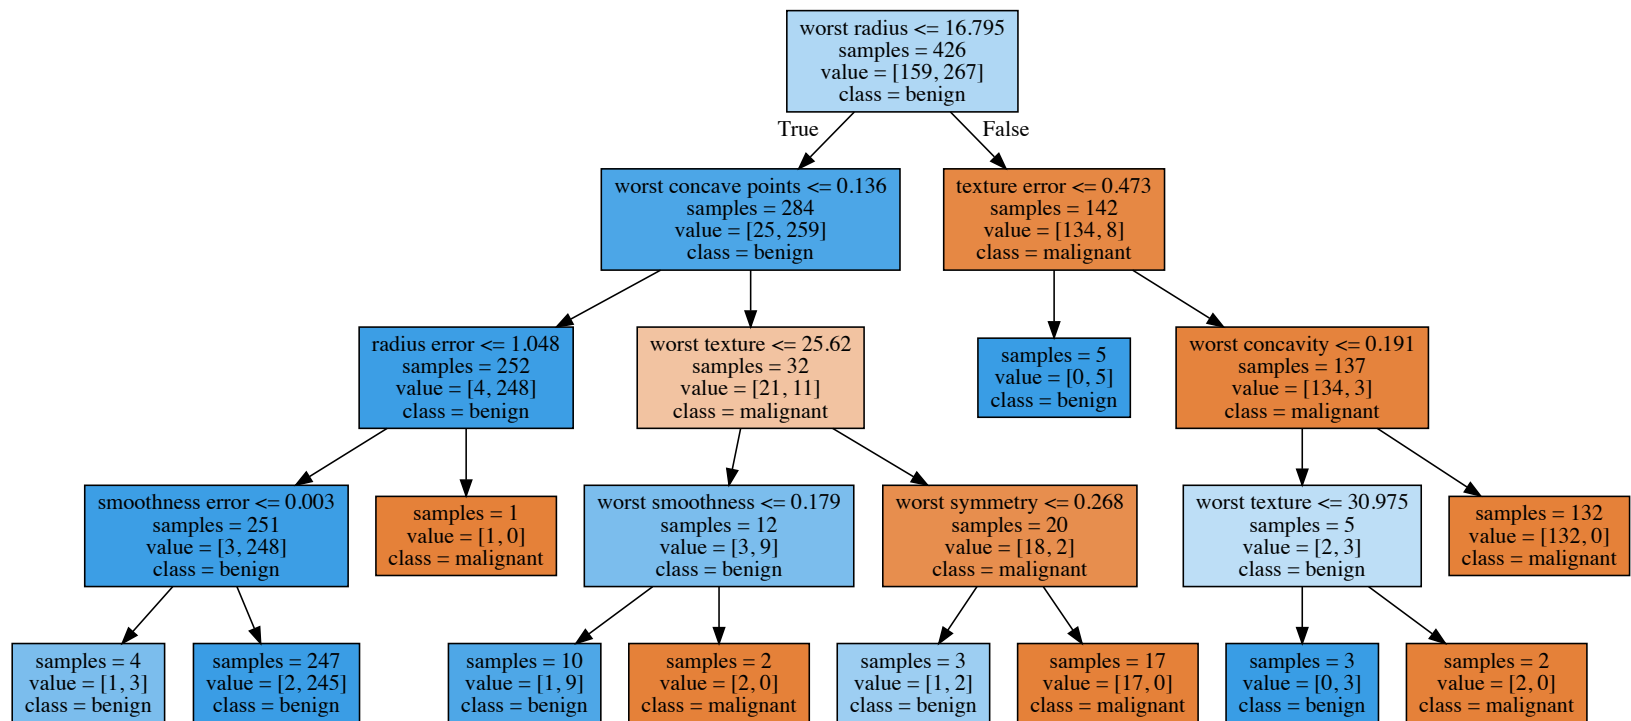
Accuracy on training set: 1.000  
Accuracy on test set: 0.937

```
tree = DecisionTreeClassifier(max_depth=5, random_state=0)
tree.fit(X_train, y_train)
```

Accuracy on training set: 0.995  
Accuracy on test set: 0.951


## Analyzing Decision Trees manually

- Visualize and find the path that most data takes



`DecisionTreeClassifier` also returns *feature importances*

- In  $[0,1]$ , sum up to 1
- High values for features selected by the algorithm, others don't contribute new information given the selected features



# Decision tree regression

- Heuristic: *Minimal quadratic distance*
- Consider splits at every data point for every variable (or halfway between)
- Dividing the data on  $X_j$  at splitpoint  $s$  leads to the following half-spaces:

$$R_1(j, s) = X : X_j \leq s \quad \text{and} \quad R_2(j, s) = X : X_j > s$$

- The best split, with predicted value  $c_i$  and actual value  $y_i$ :

$$\min_{j,s} \left( \min_{c_1} \sum_{x_i \in R_1(j,s)} (y_i - c_1)^2 + \min_{c_2} \sum_{x_i \in R_2(j,s)} (y_i - c_2)^2 \right)$$


- Assuming that the tree predicts  $y_i$  as the average of all  $x_i$  in the leaf:

$$\hat{c}_1 = \text{avg}(y_i | x_i \in R_1(j, s)) \quad \text{and} \quad \hat{c}_2 = \text{avg}(y_i | x_i \in R_2(j, s))$$

with  $x_i$  being the  $i$ -th example in the data, with target value  $y_i$


## In scikit-learn

Regression is done with `DecisionTreeRegressor`




Note that decision trees do not extrapolate well.

- The leafs return the same *mean* value no matter how far the new data point lies from the training examples.
- Example on the `ram_price` forecasting dataset







## **Strengths, weaknesses and parameters**

Pre-pruning: regularize by:

- Setting a low `max_depth`, `max_leaf_nodes`
- Setting a higher `min_samples_leaf` (default=1)

Decision trees:

- Work well with features on completely different scales, or a mix of binary and continuous features
  - Does not require normalization
- Interpretable, easily visualized
- Still tend to overfit easily. Use ensembles of trees.

# Ensemble learning

Ensembles are methods that combine multiple machine learning models to create more powerful models. Most popular are:

- **RandomForests:** Build randomized trees on random samples of the data
- **Gradient boosting machines:** Build trees iteratively, giving higher weights to the points misclassified by previous trees

In both cases, predictions are made by doing a vote over the members of the example.

**Stacking** is another technique that builds a (meta)model over the predictions of each member.

# RandomForests


Reduce overfitting by averaging out individual predictions (variance reduction)

- Take a *bootstrap sample* of your data
  - Randomly sample with replacement
  - Build a tree on each bootstrap
- Repeat `n_estimators` times
  - Higher values: more trees, more smoothing
  - Make prediction by aggregating the individual tree predictions
    - a.k.a. Bootstrap aggregating (Bagging)
- RandomForest: Randomize trees by considering only a random subset of features of size `max_features` *in each node*
  - Small `max_features` yields more different trees, more smoothing
  - Default:  $\sqrt{n\_features}$  for classification,  $\log_2(n\_features)$  for regression


Making predictions:

- Classification: soft voting (softmax)
  - Every member returns probability for each class
  - After averaging, the class with highest probability wins
- Regression:
  - Return the *mean* of all predictions


Tree 0




Tree 1




Tree 2




Tree 3



Tree 4



Random Forest



Scikit-learn algorithms:

- `RandomForestClassifier` (or `Regressor`)
- `ExtraTreesClassifier`: Grows deeper trees, faster

Most important parameters:

- `n_estimators` (higher is better, but diminishing returns)
  - Will start to underfit (bias error component increases slightly)
- `max_features` (default is typically ok)
  - Set smaller to reduce space/time requirements
- parameters of trees, e.g. `max_depth` (less effect)

`n_jobs` sets the number of parallel cores to run

`random_state` should be fixed for reproducibility

```
forest = RandomForestClassifier(n_estimators=100, random_state=0) # Vary  
forest.fit(X_train, y_train)
```

Accuracy on training set: 1.000

Accuracy on test set: 0.972




RandomForest allow another way to evaluate performance: out-of-bag (OOB) error

- While growing forest, estimate test error from training samples
- For each tree grown, 33-36% of samples are not selected in bootstrap
  - Called the 'out of bootstrap' (OOB) samples
  - Predictions are made as if they were novel test samples
  - Through book-keeping, majority vote is computed for all OOB samples from all trees
- OOB estimated test error is rather accurate in practice
  - As good as CV estimates, but can be computed on the fly (without repeated model fitting)
  - Tends to be slightly pessimistic


In scikit-learn OOB error are returned as follows:

```
oob_error = 1 - clf.oob_score_
```



## Feature importance

RandomForests provide more reliable feature importances, based on many alternative hypotheses (trees)



## **Strengths, weaknesses and parameters**

RandomForest are among most widely used algorithms:

- Don't require a lot of tuning
- Typically very accurate models
- Handles heterogeneous features well
- Implicitly selects most relevant features

Downsides:


- less interpretable, slower to train (but parallelizable)
- don't work well on high dimensional sparse data (e.g. text)

# Gradient Boosted Regression Trees (Gradient Boosting Machines)

Instead of reducing the variance of overfitted models, reduce the bias of underfitted models


- Use strong pre-pruning to build very shallow trees
  - Default `max_depth=3`
- Iteratively build new trees by increasing weights of points that were badly predicted
- Example of *additive modelling*: each tree depends on the outcome of previous trees
- Optimization: find optimal weights for all data points
  - Gradient descent (covered later) finds optimal set of weights
  - `learning_rate` controls how strongly the weights are altered in each iteration (default 0.1)
- Repeat `n_estimators` times (default 100)

Example:




After 1 iteration

- The simple decision tree divides space
- Misclassified points get higher weight (larger dots)




After 3 iterations



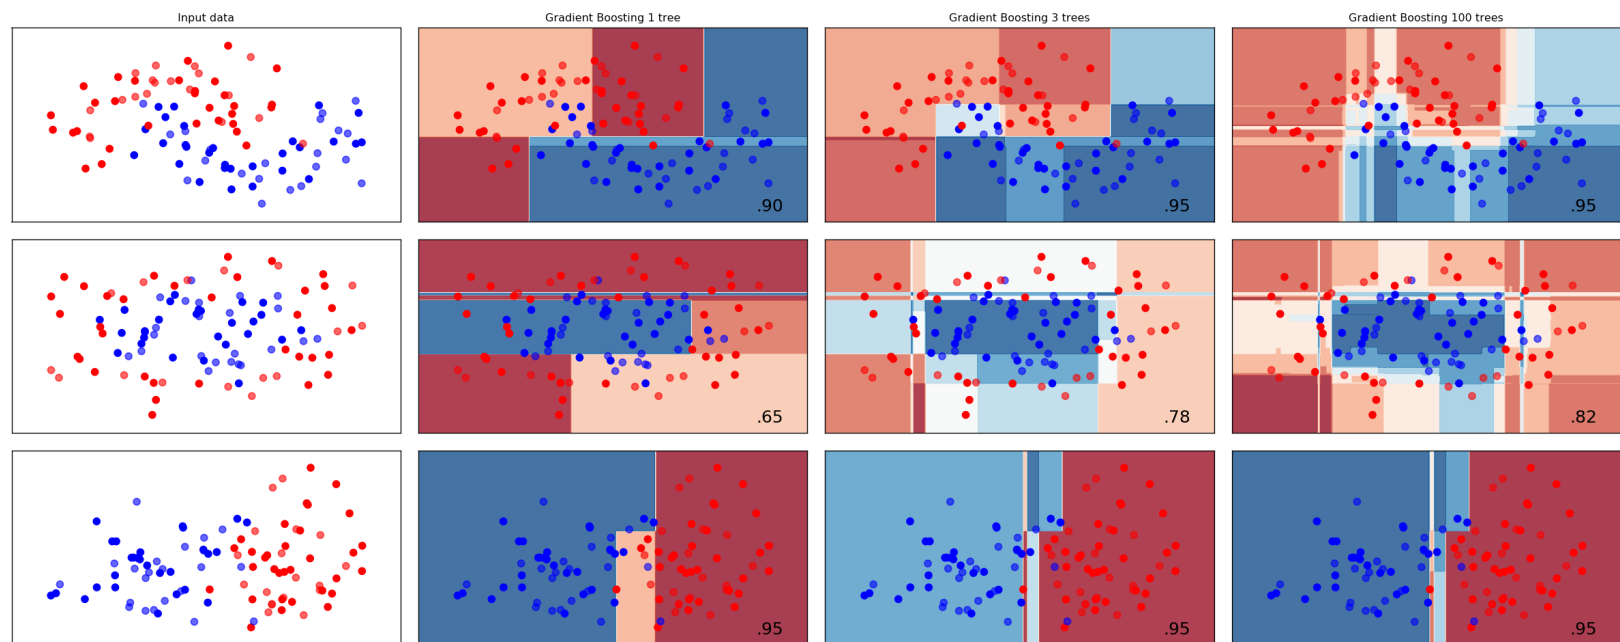


After 20 iterations



Each tree provides good predictions on part of the data, use voting for final prediction

- Soft voting for classification, mean values for regression



# Tuning

- `n_estimators`: Higher is better, but will start to overfit
- `learning_rate`: Lower rates mean more trees are needed to get more complex models
  - Main regularizer, also known as 'shrinkage'
  - Set `n_estimators` as high as possible, then tune `learning_rate`
- `max_depth`: typically kept low (<5), reduce when overfitting
- `loss`: Loss function used for gradient descent (defaults OK)
  - Classification:
    - `deviance` (default): log-likelihood loss (as in logistic regression)
    - `exponential`: exponential loss (AdaBoost algorithm)
  - Regression:
    - `ls`: Least squares (typically the best option)

```
gbrt = GradientBoostingClassifier(random_state=0)
gbrt.fit(X_train, y_train)
```

Accuracy on training set: 1.000  
Accuracy on test set: 0.958

```
gbrt = GradientBoostingClassifier(random_state=0, max_depth=1)
gbrt.fit(X_train, y_train)
```


Accuracy on training set: 0.991  
Accuracy on test set: 0.972

```
gbrt = GradientBoostingClassifier(random_state=0, learning_rate=0.01)
gbrt.fit(X_train, y_train)
```

Accuracy on training set: 0.988  
Accuracy on test set: 0.965

## Gradient boosting machines use much simpler trees

- Hence, tends to completely ignore some of the features



### *Strengths, weaknesses and parameters*

- Among the most powerful and widely used models
- Work well on heterogeneous features and different scales
- Require careful tuning, take longer to train.
- Does not work well on high-dimensional sparse data

Main hyperparameters:

- `n_estimators`: Higher is better, but will start to overfit
- `learning_rate`: Lower rates mean more trees are needed to get more complex models
  - Set `n_estimators` as high as possible, then tune `learning_rate`
- `max_depth`: typically kept low (<5), reduce when overfitting

# XGBoost

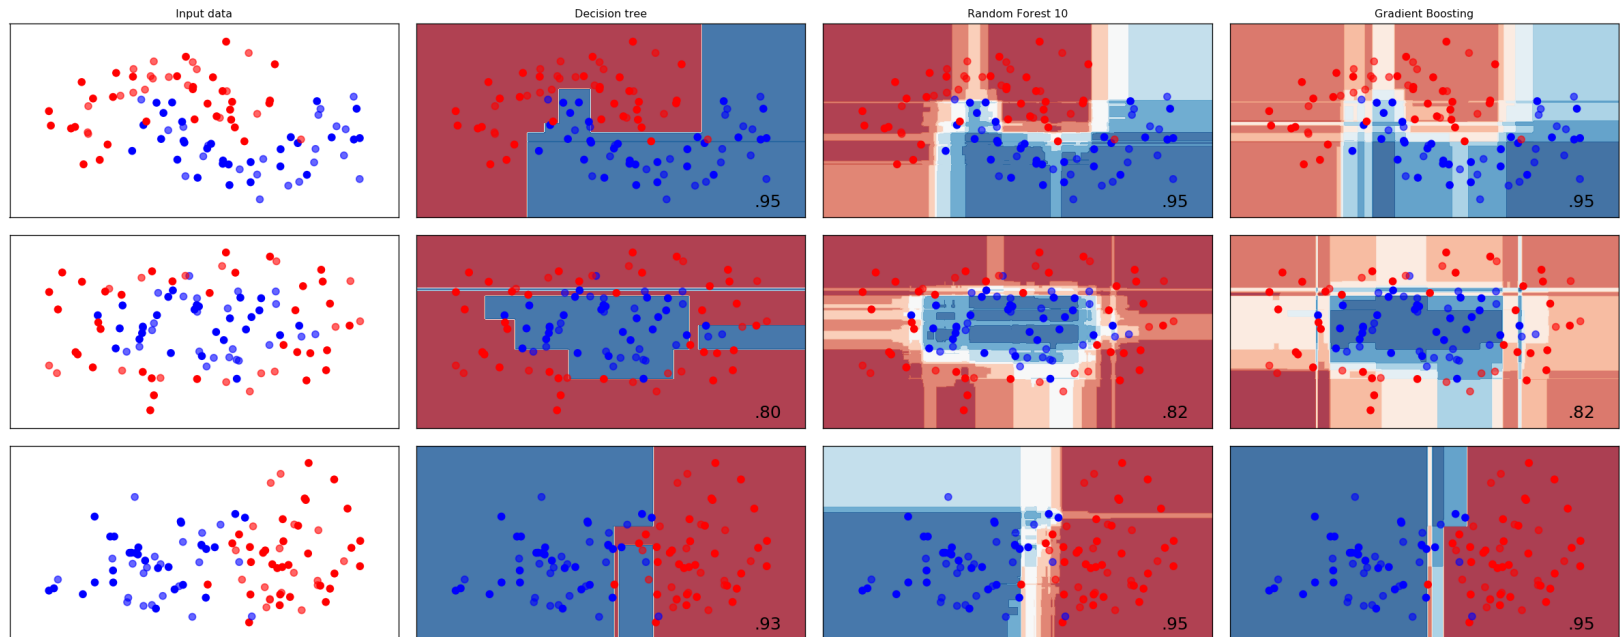
XGBoost is another python library for gradient boosting (install separately).

- The main difference lies the use of approximation techniques to make it faster.
  - Hence, you can do 10x (or 100x) more boosting iterations in same amount of time
- Sketching: Given 10000 possible splits, it will only consider 300 "good enough" splits by default
  - Controlled by the `sketch_eps` parameter (default 0.03)
- Loss function approximation with Taylor Expansion: more efficient way to evaluate splits

Further reading: [XGBoost Documentation](https://xgboost.readthedocs.io/en/latest/parameter.html#parameters-for-tree-booster)

(<https://xgboost.readthedocs.io/en/latest/parameter.html#parameters-for-tree-booster>), [Paper](http://arxiv.org/abs/1603.02754) (<http://arxiv.org/abs/1603.02754>)

# Comparison





# Summary

- Bagging / RandomForest is a variance-reduction technique
  - Build many high-variance (overfitting) models
    - Typically deep (randomized) decision trees
    - The more different the models, the better
  - Aggregation (soft voting or averaging) reduces variance
  - Parallellizes easily
- Boosting is a bias-reduction technique
  - Build many high-bias (underfitting) models
    - Typically shallow decision trees
    - Sample weights are updated to create different trees
  - Aggregation (soft voting or averaging) reduces bias
  - Doesn't parallelize easily
- You can build ensembles with other models as well
  - Especially if they show high variance or bias
- It is also possible to build *heterogeneous* ensembles
  - Models from different algorithms
  - Often a meta-classifier is trained on the predictions: Stacking